

Machine Learning With TensorFlow

X433.7-001 (2 semester units in COMPSCI)

Instructor Alexander I. Iliev, Ph.D.

Machine Learning With TensorFlow

Class 2 ...

TensorFlow and TensorBoard, Data types, Linear algebra fundamentals ...

Course Content Outline

- **Machine Learning With TensorFlow®**
 - Introduction, Python - pros and cons
 - Python modules, DL packages and scientific blocks
 - Working with the shell, IPython and the editor
 - Installing the environment with core packages
 - Writing “Hello World”
- HW1 (5pts)
- **Tensorflow and TensorBoard basics**
 - Linear algebra recap
 - Data types in Numpy and Tensorflow
 - Basic operations in Tensorflow
 - Graph models and structures with Tensorboard
- **TensorFlow operations**
 - Overloaded operators
 - Using Aliases
 - Sessions, graphs, variables, placeholders
 - Name scopes
- HW2 (5pts)
- **Data Mining and Machine Learning concepts**
 - Basic Deep Learning Models
 - Linear and Logistic Regression
 - Softmax classification
- **Neural Networks**
 - Multi-layer Neural Network
 - Gradient descent and Backpropagation
 - Object recognition with Convolutional Neural Network (CNN)
 - Activation Functions
- HW3 (5pts)

Linear algebra

- **Vector**: 1) is a **mathematical** object that has a direction and a magnitude, used to find the position of one point in space relative to another point. 2) is a **computer** object, an array of data with individual items located with a single index
- **Matrix** is a **2-dimensional (rectangular) array** of elements represented by: symbols, numbers, or expressions, all arranged in rows and columns. Matrix consist of vectors
- **Array** is an **arrangement or a series of elements** such as symbols, numbers, or expressions. Arrays can be n-dimensional, so matrix is an array with 2 dimensions
- **Tensor** is an objects describing the **linear relationship** among **scalars**, **vectors** and **other tensors**
- **Rank** of a matrix is the **maximum number of linearly independent column** (or **row**) vectors in the matrix


Linear algebra

- **Vector**

- 1) is a **mathematical** object that has a direction and a magnitude, used to find the position of one point in space relative to another point
- 2) is a **computer** object, an array of data with individual items located with a single index

Example:

mathematical meaning



computer meaning

[3 2 4 5 -6 9 2 10]

Linear algebra

- **Matrix** is a 2-dimensional (rectangular) array of elements represented by: symbols, numbers, or expressions, all arranged in rows and columns. Matrix consist of vectors

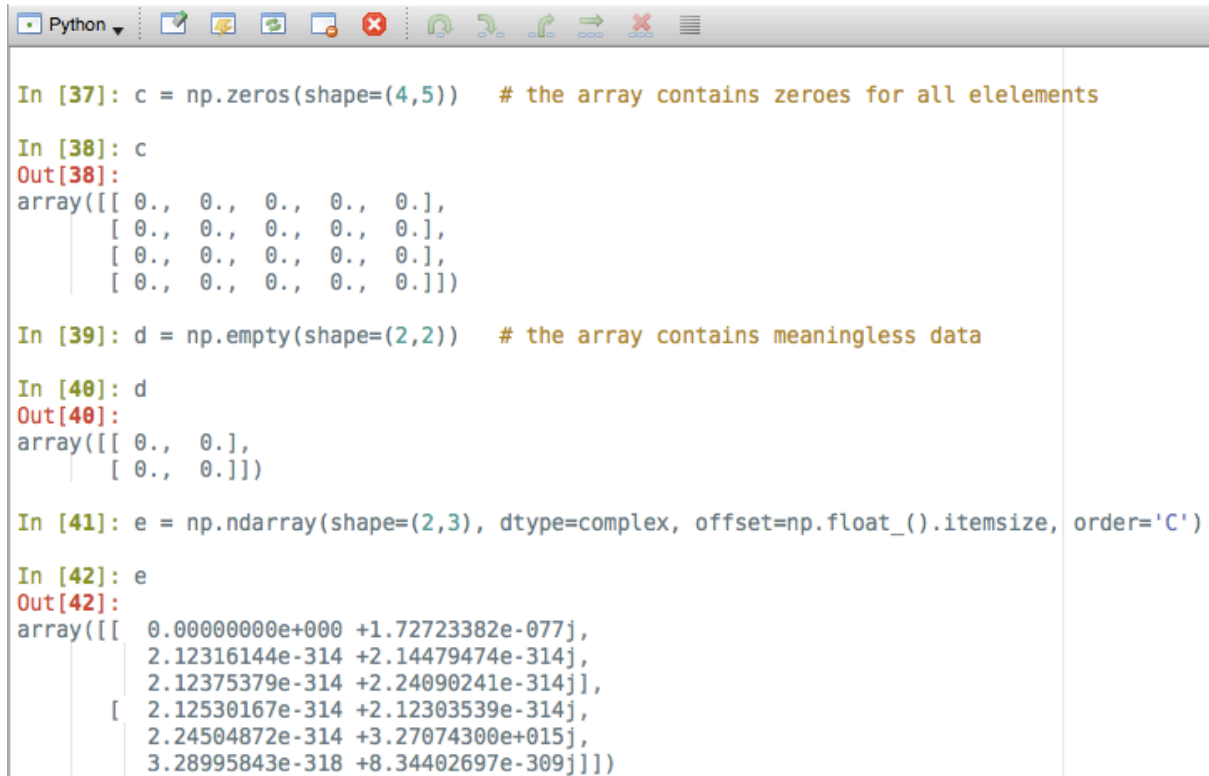
Example:

```
Python
In [1]: a = matrix([[2,-4,6],[-12,5,1],[-3,8,4]])
In [2]: a
Out[2]:
matrix([[ 2, -4,  6],
        [-12,  5,  1],
        [-3,  8,  4]])
In [3]: b = matrix([[2],[8],[-4]])
In [4]: b
Out[4]:
matrix([[ 2],
        [ 8],
        [-4]])
In [5]: a*b
Out[5]:
matrix([[ -52],
        [ 12],
        [ 42]])
```

Linear algebra

- **Array** can be n-dimensional, but matrix is an array with 2 dimensions

Example: `numpy.array` is a function that returns a `numpy.ndarray` and there for convenience



```
Python
In [37]: c = np.zeros(shape=(4,5)) # the array contains zeroes for all elements

In [38]: c
Out[38]:
array([[ 0.,  0.,  0.,  0.,  0.],
       [ 0.,  0.,  0.,  0.,  0.],
       [ 0.,  0.,  0.,  0.,  0.],
       [ 0.,  0.,  0.,  0.,  0.]])

In [39]: d = np.empty(shape=(2,2)) # the array contains meaningless data

In [40]: d
Out[40]:
array([[ 0.,  0.],
       [ 0.,  0.]])

In [41]: e = np.ndarray(shape=(2,3), dtype=complex, offset=np.float_().itemsize, order='C')

In [42]: e
Out[42]:
array([[ 0.00000000e+000 +1.72723382e-077j,
        2.12316144e-314 +2.14479474e-314j,
        2.12375379e-314 +2.24090241e-314j],
       [ 2.12530167e-314 +2.12303539e-314j,
        2.24504872e-314 +3.27074300e+015j,
        3.28995843e-318 +8.34402697e-309j]])
```

Linear algebra

- NumPy **array**

Example:

```
Python
In [23]: a = np.array([[12, 34, 41], [54, 62, 18], [72, 84, 96]], np.int16)

In [24]: a
Out[24]:
array([[12, 34, 41],
       [54, 62, 18],
       [72, 84, 96]], dtype=int16)

In [25]: a.size
Out[25]: 9

In [26]: a.shape
Out[26]: (3, 3)

In [27]: type(a)
Out[27]: numpy.ndarray

In [28]: a.dtype
Out[28]: dtype('int16')

In [29]: a[2,2] # this is how we index a particular elemnt in the array (#9)
Out[29]: 96

In [30]: b = a[0,:]

In [31]: b
Out[31]: array([12, 34, 41], dtype=int16)

In [32]: b.shape
Out[32]: (3,)

In [33]: b[2] = 88 # this is how we reassign another value to a member in the array

In [34]: a[2,2] = 99 # the change above also affects the original array

In [35]: a
Out[35]:
array([[12, 34, 88],
       [54, 62, 18],
       [72, 84, 99]], dtype=int16)

In [36]: b
Out[36]: array([12, 34, 88], dtype=int16)
```


Linear algebra

- Difference between a `numpy.matrix` and 2D `numpy.ndarray`
 - basic operations such as `multiplications` and `transpose` are included in NumPy for both `matrix` and `ndarray` types
 - `numpy.matrix` has a `proper interface` for matrix operations than `numpy.ndarray`
 - however, the `numpy.matrix` class `does not add anything` that cannot be achieved by using 2D `numpy.ndarray` objects
 - the implementation of this class `resembles` the one in Matlab
 - `scipy.linalg` operations can be used `just as good to 2D numpy.ndarray` objects as well as to `numpy.matrix`

Linear algebra

- Difference between a `numpy.matrix` and 2D `numpy.ndarray`

the `I` and `T` class members serve as shortcuts for **inverse** and **transpose** respectively

the `numpy.matrix` class does not add anything that cannot be achieved by using 2D `numpy.ndarray` objects

```
Python
In [45]: # Matrix vs Array:

In [46]: # 1. Matrix:

In [47]: import numpy as npy

In [48]: a = npy.mat('[12 34 41;52 64 72]') # create matrix 'a'

In [49]: a
Out[49]:
matrix([[12, 34, 41],
        [52, 64, 72]])

In [50]: type(a)
Out[50]: numpy.matrixlib.defmatrix.matrix

In [51]: a.I # inverse of matrix 'a'
Out[51]:
matrix([[ -0.05885099,  0.03258602],
        [ 0.01490374, -0.00181295],
        [ 0.02925572, -0.00803395]])

In [52]: b = npy.mat('[84 92]') # create matrix 'b'

In [53]: b
Out[53]: matrix([[84, 92]])

In [54]: type(b)
Out[54]: numpy.matrixlib.defmatrix.matrix

In [55]: b.T # transpose of matrix 'b'
Out[55]:
matrix([[84],
        [92]])

In [56]: a.T*b.T # multiplication of two matrices
Out[56]:
matrix([[ 5792],
        [ 8744],
        [10068]])
```

Linear algebra

- Difference between a `numpy.matrix` and 2D `numpy.ndarray`

`scipy.linalg` operations can be used just as good to 2D `numpy.ndarray` objects as well as to `numpy.matrix`

Note:

`numpy.mat` and `numpy.matrix` are the same. Try, using 'id'

```
Python
In [57]: # Matrix vs Array:
In [58]: # 2. Array:
In [59]: import numpy as npy
In [60]: from scipy import linalg
In [61]: c = npy.array([[12,34,41],[52,64,72],[84,91,98]]) # create array 'c'
In [62]: c
Out[62]:
array([[12, 34, 41],
       [52, 64, 72],
       [84, 91, 98]])
In [63]: type(c)
Out[63]: numpy.ndarray
In [64]: linalg.inv(c) # calculate the inverse of a matrix
Out[64]:
array([[ -0.10752688,  0.15322581, -0.06758833],
       [ 0.3655914 , -0.87096774,  0.48694316],
       [-0.24731183,  0.67741935, -0.38402458]])
In [65]: d2 = npy.array([[2,12,28]]) # create 2D array
In [66]: d2
Out[66]: array([[ 2, 12, 28]])
In [67]: type(d2)
Out[67]: numpy.ndarray
```

Linear algebra

- Difference between a `numpy.matrix` and 2D `numpy.ndarray`

`scipy.linalg` operations
can be used just as
good to 2D
`numpy.ndarray`
objects as well as to
`numpy.matrix`

```
Python
In [68]: c*d2 # this is not matrix multiplication
Out[68]:
array([[ 24,  408, 1148],
       [ 104,  768, 2016],
       [ 168, 1092, 2744]])

In [69]: c.dot(d2.T) # this is matrix multiplication using dot product function
Out[69]:
array([[1580],
       [2888],
       [4004]])

In [70]: d1 = npy.array([2,12,28]) # this is 1D array

In [71]: d1
Out[71]: array([ 2, 12, 28])

In [72]: type(d1)
Out[72]: numpy.ndarray

In [73]: d1.T # this is not a matrix transpose
Out[73]: array([ 2, 12, 28])

In [74]: c.dot(d1) # dot product of two arrays
Out[74]: array([1580, 2888, 4004])
```

Linear algebra

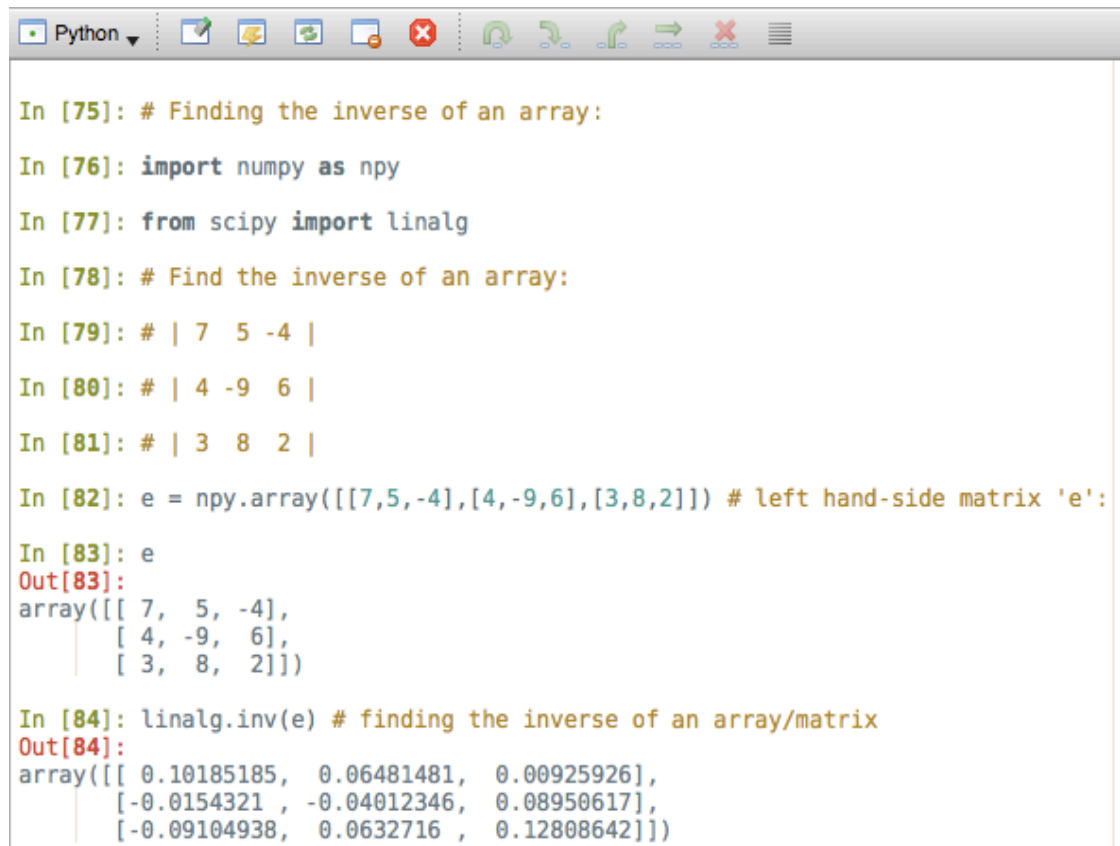
- Finding the inverse of an array (matrix)

matrix 'e' has its **inverse** matrix 'f' such that $e*f=I \rightarrow I$ is the **identity** matrix that has main diagonal with ones

we can then say that: $f=e^{-1}$

the **matrix inverse** of the NumPy array 'e' is obtained in **two ways**:

- using the Scipy `linalg.inv`, or
- using `e.I` when 'e' is a **matrix** so cast it like this:
`numpy.mat(e).I`




```
Python
In [75]: # Finding the inverse of an array:
In [76]: import numpy as npy
In [77]: from scipy import linalg
In [78]: # Find the inverse of an array:
In [79]: # | 7  5 -4 |
In [80]: # | 4 -9  6 |
In [81]: # | 3  8  2 |
In [82]: e = npy.array([[7,5,-4],[4,-9,6],[3,8,2]]) # left hand-side matrix 'e':
In [83]: e
Out[83]:
array([[ 7,  5, -4],
       [ 4, -9,  6],
       [ 3,  8,  2]])
In [84]: linalg.inv(e) # finding the inverse of an array/matrix
Out[84]:
array([[ 0.10185185,  0.06481481,  0.00925926],
       [-0.0154321 , -0.04012346,  0.08950617],
       [-0.09104938,  0.0632716 ,  0.12808642]])
```

Linear algebra

- **Tensor** is an objects describing the **linear relationship** among **scalars**, **vectors** and **other tensors**

Example:

a 2nd order **tensor** of a
3-dimensional space
represent the **matrix**:



$$\sigma = [T(e_1)T(e_2)T(e_3)] \quad \text{or} \quad \sigma = \begin{bmatrix} \sigma_{11} & \sigma_{12} & \sigma_{13} \\ \sigma_{21} & \sigma_{22} & \sigma_{23} \\ \sigma_{31} & \sigma_{32} & \sigma_{33} \end{bmatrix}$$

where, the columns are the forces e_n depicted
on the 3 faces of the cube (e_1, e_2, e_3)

Linear algebra

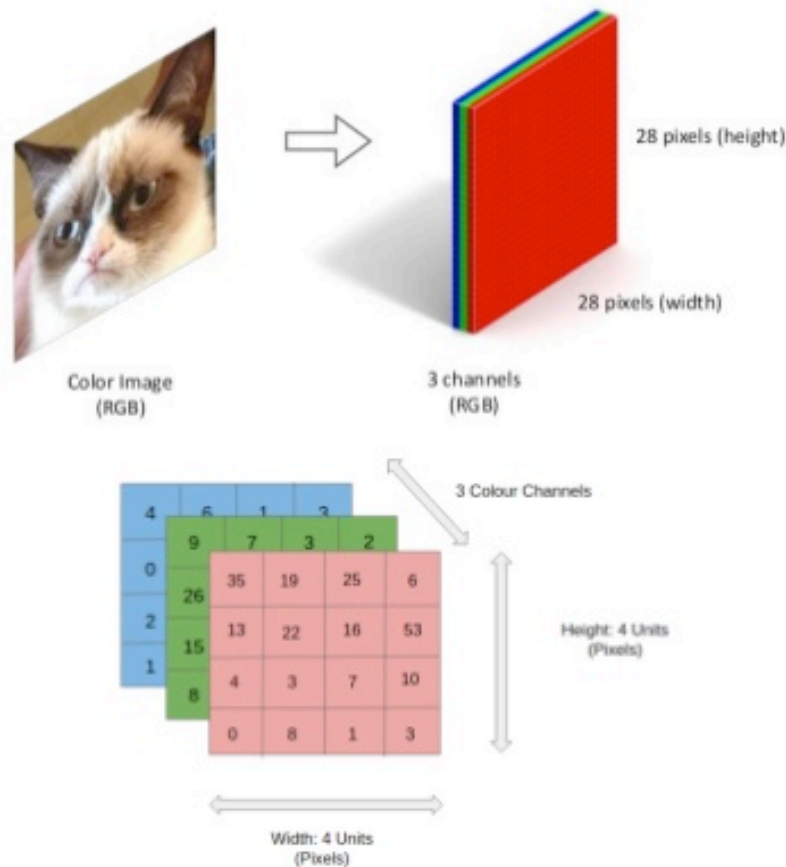
- **Tensor** is an objects describing the **linear relationship** among **scalars**, **vectors** and **other tensors**
 - A 0th order tensor can be represented by a **scalar**
 - A 1st order tensor can be represented by an **array (vector)**
 - A 2nd order tensor can be represented by a **matrix**
 - A 3rd order tensor can be represented as a **3-dimensional array** of numbers
 - However **tensor** represents more than just an arrangement of components:
 - **tensor** shows how the array transforms upon a change of its basis
 - **tensor** is an **ndarray** satisfying a particular transformation law

Linear algebra

- Rank of a matrix is the maximum number of linearly independent column (or row) vectors in the matrix, denoted as: $rk(A)$ or $rank(A)$
 - The rank R of a tensor is independent of the number of dimensions N of the underlying space
 - Rank-0 is a scalar $N^0 = 1$
 - Rank-1 is an array (vector) $N^1 = N$
 - Rank-2 is a matrix $N^2 = N \times N$ aka dyad, dyadic
 - Rank-3 is a 3-darray $N^3 = N \times N \times N$ aka triad
 - Rank-4 is a 4-darray $N^4 = N \times N \times N \times N$ aka tetrad
 - etc

Linear algebra

color image is 3rd-order tensor



Linear algebra

- **Rank** of a matrix is the maximum number of linearly independent column (or row) vectors in the matrix, denoted as: $rk(A)$ or $rank(A)$

– How to find the rank:

- They are generally: 0,1,2 and 3:

Rank(A) = 0 when matrix is null

Rank(A) = 1 when every sub-matrix of A is singular or $\det(A_n) = 0$

Rank(A) = 2 when A is singular or $|A| = 0$, **and** at least one of its sub-matrix is $|A_1| \neq 0$

Rank(A) = 3 when A is non-singular or $|A| \neq 0$

Example:

$$A = \begin{bmatrix} -1 & -1 & 0 \\ 4 & 2 & 2 \\ 3 & 1 & 2 \end{bmatrix}, \quad \begin{aligned} 1) \det(A) &= (-1)(2 \cdot 2 - 2 \cdot 1) - (-1)(4 \cdot 2 - 2 \cdot 3) + (0)(4 \cdot 1 - 2 \cdot 3) = -2 + 2 - 0 = 0 \\ &\Rightarrow |A| = 0 \text{ is singular} \\ 2) \det(A_1) &= 4 \cdot 1 - 2 \cdot 3 = 4 - 6 = -2 \quad \text{--} > A_1 \text{ is a sub-matrix of } A \\ &\Rightarrow |A_1| \neq 0 \text{ and is non-singular} \end{aligned}$$

Therefore Rank(A) = 2

Linear algebra

- **Rank** of a matrix is the maximum number of linearly independent column (or row) vectors in the matrix, denoted as: $rk(A)$ or $rank(A)$

– How to find the rank:

- They are generally: 0,1,2 and 3:

Rank(A) = 0 when matrix is null

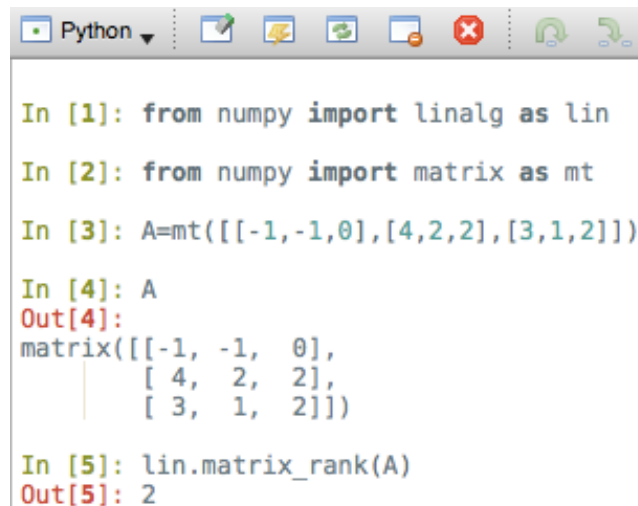
Rank(A) = 1 when every sub-matrix of A is singular or $\det(A_n) = 0$

Rank(A) = 2 when A is singular or $|A| = 0$, **and** at least one of its sub-matrix is $|A_1| \neq 0$

Rank(A) = 3 when A is non-singular or $|A| \neq 0$

Example:

$$A = \begin{bmatrix} -1 & -1 & 0 \\ 4 & 2 & 2 \\ 3 & 1 & 2 \end{bmatrix}$$



```
Python
In [1]: from numpy import linalg as lin
In [2]: from numpy import matrix as mt
In [3]: A=mt([[-1,-1,0],[4,2,2],[3,1,2]])
In [4]: A
Out[4]:
matrix([[-1, -1,  0],
        [ 4,  2,  2],
        [ 3,  1,  2]])
In [5]: lin.matrix_rank(A)
Out[5]: 2
```

NumPy data type objects

- Data type objects
 - NumPy supports much larger variety of types than what the standard Python implementation does:

Number type	Data type	Description
Booleans	bool, bool8, bool_	Boolean (True or False) stored as a byte – 8 bits
Integers	byte	compatible: C char – 8 bits
	short	compatible: C short – 16 bits
	int, int0, int_	Default integer type (same as C long; normally either int32 or int64) – 64 bits
	longlong	compatible: C long long – 64 bits
	intc	Identical to C int – 32 bits
	intp	Integer used for indexing (same as C size_t) – 64 bits
	int8	Byte (-128 to 127) – 8 bits
	int16	Integer (-32768 to 32767) – 16 bits
	int32	Integer (-2147483648 to 2147483647) – 32 bits
	int64	Integer (-9223372036854775808 to 9223372036854775807) – 64 bits
Unsigned integers	uint, uint0	Python int compatible, unsigned – 64 bits
	ubyte	compatible: C unsigned char, unsigned – 8 bits
	ushort	compatible: C unsigned short, unsigned – 16 bits
	ulonglong	compatible: C long long, unsigned – 64 bits
	uintp	large enough to fit a pointer – 64 bits
	uintc	compatible: C unsigned int – 32 bits
	uint8	Unsigned integer (0 to 255) – 8 bits
	uint16	Unsigned integer (0 to 65535) – 16 bits
	uint32	Unsigned integer (0 to 4294967295) – 32 bits
	uint64	Unsigned integer (0 to 18446744073709551615) – 64 bits

NumPy data type objects

- Data type objects
 - NumPy supports much larger variety of types than what the standard Python implementation does:

Number type	Data type	Description
Floating-point numbers	half	compatible: C short – 16 bits
	single	compatible: C float – 32 bits
	double	compatible: C double – 64 bits
	longfloat	compatible: C long float – 128 bits
	float_	Shorthand for float64 – 64 bits
	float16	Half precision float: sign bit, 5 bits exponent, 10 bits mantissa
	float32	Single precision float: sign bit, 8 bits exponent, 23 bits mantissa
	float64	Double precision float: sign bit, 11 bits exponent, 52 bits mantissa
	float128	128 bits
Complex floating-point numbers	csingle	64 bits
	complex, complex_	Shorthand for complex128 – 128 bits
	complex64	Complex number, represented by two 32-bit floats (real and imaginary components)
	complex128	Complex number, represented by two 64-bit floats (real and imaginary components)
	complex256	two 256 bit floats

- To check how many bits each type occupies, use one of these notations:
 - 1) `(np.dtype(np.<type>).itemsize)*8`
 - 2) `np.<type>().itemsize*8`

NumPy data type objects

- Data type objects
 - the difference between **signed** and **unsigned** integers and long type variables is:
 - the **signed** and **unsigned** types are of the **same size**
 - the **signed** can represent **equal amount of values around the '0'** thus representing equal amount of positive and negative numbers
 - the **unsigned can not represent any negative numbers**, but can represent double the amount of total positive numbers as compared to the signed type
 - for 32-bit int we have:
 - int**: -2147483648 to 2147483647
 - uint**: 0 to 4294967295
 - for 64-bit long we have:
 - long**: -9223372036854775808 to 9223372036854775807
 - ulong**: 0 to 18446744073709551615

TensorFlow

- Below is the full list of data types available in TensorFlow:

```
29 @tf_export("DType")
30 class DType(object):
31     """Represents the type of the elements in a `Tensor`.
32
33     The following `DType` objects are defined:
34
35     * `tf.float16`: 16-bit half-precision floating-point.
36     * `tf.float32`: 32-bit single-precision floating-point.
37     * `tf.float64`: 64-bit double-precision floating-point.
38     * `tf.bfloat16`: 16-bit truncated floating-point.
39     * `tf.complex64`: 64-bit single-precision complex.
40     * `tf.complex128`: 128-bit double-precision complex.
41     * `tf.int8`: 8-bit signed integer.
42     * `tf.uint8`: 8-bit unsigned integer.
43     * `tf.uint16`: 16-bit unsigned integer.
44     * `tf.uint32`: 32-bit unsigned integer.
45     * `tf.uint64`: 64-bit unsigned integer.
46     * `tf.int16`: 16-bit signed integer.
47     * `tf.int32`: 32-bit signed integer.
48     * `tf.int64`: 64-bit signed integer.
49     * `tf.bool`: Boolean.
50     * `tf.string`: String.
51     * `tf.qint8`: Quantized 8-bit signed integer.
52     * `tf.quint8`: Quantized 8-bit unsigned integer.
53     * `tf.qint16`: Quantized 16-bit signed integer.
54     * `tf.quint16`: Quantized 16-bit unsigned integer.
55     * `tf.qint32`: Quantized 32-bit signed integer.
56     * `tf.resource`: Handle to a mutable resource.
57     * `tf.variant`: Values of arbitrary types.
```

TensorFlow

- Examples:

```
Python
In [1]: import tensorflow as tf
In [2]: import numpy as np
In [3]: sess = tf.Session()
In [4]: a = 92
In [5]: type(a)
Out[5]: int
In [6]: a
Out[6]: 92
In [7]: a = tf.convert_to_tensor(a, dtype = tf.float16)
In [8]: type(a)
Out[8]: tensorflow.python.framework.ops.Tensor
```

... try it in class

```
Python
In [9]: a
Out[9]: <tf.Tensor 'Const:0' shape=() dtype=float16>
In [10]: sess.run(a)
Out[10]: 92.0
In [11]: type(sess.run(a))
Out[11]: numpy.float16
In [12]: a = sess.run(a)
In [13]: type(a)
Out[13]: numpy.float16
In [14]: a
Out[14]: 92.0
In [15]: a = tf.cast(a, tf.float32)
In [16]: type(a)
Out[16]: tensorflow.python.framework.ops.Tensor
In [22]: a = tf.cast(a, np.float64)
In [23]: type(a)
Out[23]: tensorflow.python.framework.ops.Tensor
In [24]: type(sess.run(a))
Out[24]: numpy.float64
In [27]: a = np.int16(sess.run(a))
In [28]: type(a)
Out[28]: numpy.int16
In [29]: a
Out[29]: 92
```


TensorFlow

- Few extra notes:
 - **Graphs** save computation time
 - **Graphs** break computation into small pieces to facilitates auto-differentiation
 - **Graphs** handle distributed computation, that is they spread the work across multiple CPUs, GPUs, or devices
 - Many machine learning models are taught and visualized as directed **graphs**
 - **Nodes** in the graph are called **ops** (short for **operations**)
 - An **op** takes zero or more Tensors, performs some computation, and produces zero or more Tensors
 - **Sessions** gives us the **environment** to perform **operations** on our **tensor data**

TensorFlow

- Few extra notes on types:
 - Using Python types to specify Tensor objects is quick and easy, and it is useful for prototyping ideas
 - However, there is an important and unfortunate downside to doing it this way:
 - TensorFlow has a plethora of data types, but basic Python types lack the ability to explicitly state what kind of data type we'd like to use!
 - Instead, TensorFlow has to infer which data type it was meant
 - TensorFlow is tightly integrated with NumPy, the scientific computing package designed for manipulating ndarrays

TensorFlow

- Few extra notes on types:
 1. In **TensorFlow**, all data passed from node to node are **Tensor objects**
 2. **TensorFlow Operations** are able to look at standard **Python types**, such as integers and strings, and automatically **convert them into tensors**
 3. There are a **variety of ways to create Tensor objects** manually (that is, without reading it in from an external data source)
 4. ... so let's see them next: ...

TensorFlow

- Few extra notes on types:
 - TensorFlow can take in Python numbers, booleans, strings, or lists of any of the above
 - **Single values** will be converted to a **0-D Tensor** (or **scalar**)
 - **Lists of values** will be converted to a **1-D Tensor** (vector)
 - **Lists of lists** of values will be converted to a **2-D Tensor** (matrix), and so on

TensorFlow

- Here is a small chart showcasing this:

```
t_0 = 50                                # Treated as 0-D Tensor,  
or "scalar"  
  
t_1 = [b"apple", b"peach", b"grape"] # Treated as 1-D Tensor,  
or "vector"  
  
t_2 = [[True, False, False],           # Treated as 2-D Tensor,  
or "matrix"  
        [False, False, True],  
        [False, True, False]]  
  
t_3 = [[ [0, 0], [0, 1], [0, 2] ],      # Treated as 3-D Tensor  
        [ [1, 0], [1, 1], [1, 2] ],  
        [ [2, 0], [2, 1], [2, 2] ]]  
...
```

TensorFlow

- Few extra notes on types (cont.):
 - TensorFlow's data types are based on those from NumPy
 - **Tensors are just a superset of matrices!**
 - In fact, the statement `np.int32 == tf.int32` returns **True!**
 - Any NumPy array can be passed into any TensorFlow Op, and the beauty is that you can easily specify the data type you need with minimal effort

The String data types problem:

- For numeric and boolean types, TensorFlow and NumPy dtypes match perfectly well
- However, **tf.string does not have an exact match in NumPy** due to the way NumPy handles strings
- That said, TensorFlow can import string arrays from NumPy perfectly fine - **just don't specify a dtype in NumPy!**

TensorFlow

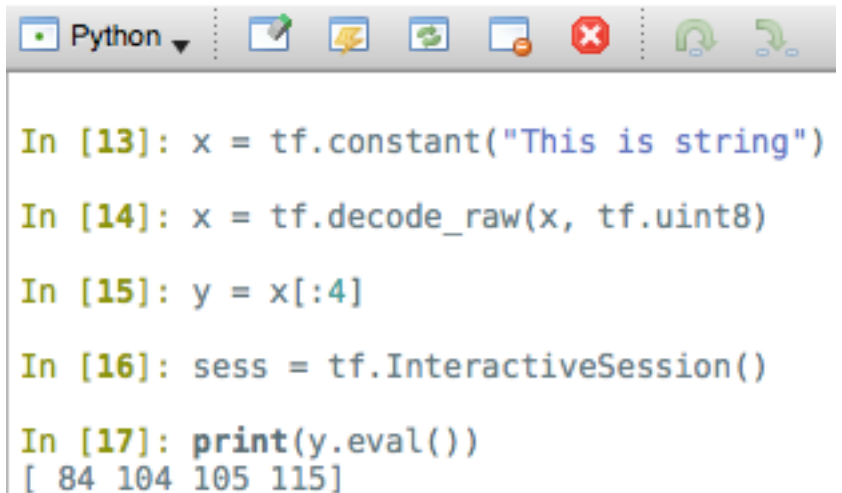
- Unlike Python, where a string can be treated as a list of characters, TensorFlow's **tf.strings** are indivisible values

Example:

'x' is a Tensor with shape (2,) and each element inside is a variable length string

```
x = tf.constant(["This is a string", "This is another string"])
```

- TensorFlow provides the **tf.decode_raw** operator that takes **tf.string** tensor as input, but can decode the string into any other primitive data type
- To interpret the string as a tensor of characters, you can do:



```
Python | [Python icon] [Run icon] [Save icon] [Refresh icon] [Close icon] [Help icon] [Undo icon] [Redo icon]

In [13]: x = tf.constant("This is string")
In [14]: x = tf.decode_raw(x, tf.uint8)
In [15]: y = x[:4]
In [16]: sess = tf.InteractiveSession()
In [17]: print(y.eval())
[ 84 104 105 115]
```

TensorFlow

- Few extra notes on types (cont.):
 - You can use the functionality of the **numpy** library both before and after running your graph, as the tensors returned from **Session.run** are NumPy arrays
 - Here's an example of how to create NumPy arrays, mirroring the previous example:

```
import numpy as np  # Don't forget to import NumPy!

# 0-D Tensor with 32-bit integer data type
t_0 = np.array(50, dtype=np.int32)

# 1-D Tensor with byte string data type
# Note: don't explicitly specify dtype when using strings in
# NumPy
t_1 = np.array([b"apple", b"peach", b"grape"])

# 1-D Tensor with boolean data type
t_2 = np.array([[True, False, False],
                [False, False, True],
```


TensorFlow

- Few extra notes on types (cont.):
 - You can use the functionality of the **numpy** library both before and after running your graph, as the tensors returned from **Session.run** are NumPy arrays
 - Here's an example of how to create NumPy arrays, mirroring the previous example:

```
[False, True, False]],  
dtype=np.bool)  
  
# 3-D Tensor with 64-bit integer data type  
t_3 = np.array([[ [0, 0], [0, 1], [0, 2] ],  
                 [ [1, 0], [1, 1], [1, 2] ],  
                 [ [2, 0], [2, 1], [2, 2] ]],  
               dtype=np.int64)  
  
...
```

Tensor Shape

- The **shape in TensorFlow** terminology **describes** both:
 - the **number of dimensions** in a tensor as well as
 - the **length of each dimension**
- **Tensor shapes** can either be **Python lists** or **tuples** containing an ordered set of integers:
 - there are as many numbers in the list as there are dimensions, and each number describes the length of its corresponding dimension

Example:

the **list [2, 3]** describes the shape of a **2-D tensor of length 2 in its first dimension and length 3 in its second dimension**

* Note that either: - tuples (wrapped with parentheses `()`) or
- lists (wrapped with brackets `[]`) can be used to define shapes

Tensor Shape

- Some examples:

```
# Shapes that specify a 0-D Tensor (scalar)
# e.g. any single number: 7, 1, 3, 4, etc.
s_0_list = []
s_0_tuple = ()

# Shape that describes a vector of length 3
# e.g. [1, 2, 3]
s_1 = [3]
```

```
# Shape that describes a 3-by-2 matrix
# e.g. [[1, 2],
#       [3, 4],
#       [5, 6]]
s_2 = (3, 2)
```

Tensor Shape

- In addition to being able to specify fixed lengths to each dimension, **we are also able assign a flexible length** by **passing in None** as a dimension's value
- This will tell TensorFlow to **allow a tensor of any shape**
- That is, a tensor with **any amount of dimensions** and **any length for each dimension**:

```
# Shape for a vector of any length:  
s_1_flex = [None]  
  
# Shape for a matrix that is any amount of rows tall, and 3  
columns wide:  
s_2_flex = (None, 3)  
  
# Shape of a 3-D Tensor with length 2 in its first dimension,  
and variable-  
# length in its second and third dimensions:  
s_3_flex = [2, None, None]  
  
# Shape that could be any Tensor  
s_any = None
```

Tensor Shape

- If we ever need to figure out the shape of a tensor in the middle of our graph, we can use the `tf.shape` Op
- It simply takes in the Tensor object we'd like to find the shape for, and returns it as an int32 vector:

```
import tensorflow as tf

# ...create some sort of mystery tensor

# Find the shape of the mystery tensor
shape = tf.shape(mystery_tensor, name="mystery_shape")
```

** Remember that `tf.shape`, like any other Operation, doesn't run until it is executed inside of a Session*

TensorFlow

- Examples:

```
1 ## Tensor Types
2 import tensorflow as tf
3 import numpy as np
4
5 # Define a 2x2 matrix in 3 different ways
6 m1 = [[1.0, 2.0], [3.0, 4.0]] # <class 'list'>
7 m2 = np.array([[1.0, 2.0], [3.0, 4.0]], dtype=np.float32) # <class 'numpy.ndarray'>
8 m3 = tf.constant([[1.0, 2.0], [3.0, 4.0]]) # <class 'tensorflow'>
9 print(type(m1))
10 print(type(m2))
11 print(type(m3))
12
13 # Create tensor objects out of various types
14 t1 = tf.convert_to_tensor(m1, dtype=tf.float32) # <class 'tensorflow'>
15 t2 = tf.convert_to_tensor(m2, dtype=tf.float32) # <class 'tensorflow'>
16 t3 = tf.convert_to_tensor(m3, dtype=tf.float32) # <class 'tensorflow'>
17 print(type(t1))
18 print(type(t2))
19 print(type(t3))
```

```
Python
In [9]: m3.op.name
Out[9]: 'Const'
```

... try it in class

```
Python
In [1]: (executing cell "Tensor Types" (line 2 of "types.py"))
<class 'list'>
<class 'numpy.ndarray'>
<class 'tensorflow.python.framework.ops.Tensor'>
<class 'tensorflow.python.framework.ops.Tensor'>
<class 'tensorflow.python.framework.ops.Tensor'>
<class 'tensorflow.python.framework.ops.Tensor'>
```

TensorFlow

- Before you can use a variable, it must be initialized
- Most high-level frameworks such as [tf.contrib.slim](#), [tf.estimator.Estimator](#) and [Keras](#) automatically initialize variables
- If you are explicitly creating your own graphs and sessions, you **must explicitly initialize the variables**
- To initialize all trainable variables in one go, **before training starts**, call:
`session.run(tf.global_variables_initializer())`
- To ask which variables have still not been initialized:
`session.run(my_variable.initializer)`

TensorFlow

- Examples:

```
1  ## A simple matrix and the InteractiveSession:
2  import tensorflow as tf
3  sess = tf.InteractiveSession()
4
5  matrix = tf.constant([[5., 6.]])
6  negMatrix = tf.negative(matrix) → [[-5. -6.]]
7
8  result = negMatrix.eval()
9  print(result)
10 sess.close()
```

- **tf.InteractiveSession:**

- The only difference with a regular `Session` is that an `InteractiveSession` installs itself as the default session on construction.

- For example:

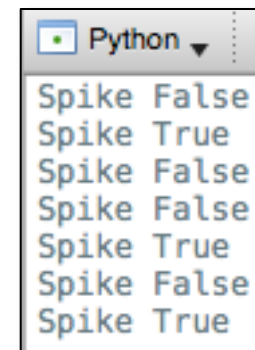
```
sess = tf.InteractiveSession()
a = tf.constant(5.0)
b = tf.constant(6.0)
c = a * b
# We can just use 'c.eval()' without passing 'sess'
print(c.eval())
sess.close()
```

... try it in class

TensorFlow

- Examples:

```
1  ## Detecting Spikes:
2  import tensorflow as tf
3  sess = tf.InteractiveSession()
4
5  # Create a boolean variable called `spike` to detect sudden a sudden increase in
6  # a series of numbers. Since all variables must be initialized, initialize the
7  # variable by calling `run()` on its `initializer`:
8  vector = [3., -2., 8., -4., 0.2, 2.3, 7.5, 14.8]
9  spike = tf.Variable(False)
10
11 # Initiallizing our variable in one of two ways (choose one):
12 # sess.run(spike.initializer)
13 spike.initializer.run()
14
15 # Loop through the data and update the spike variable when there is a
16 # significant increase:
17 for i in range(1, len(vector)):
18     if vector[i] - vector[i-1] > 5:
19         updater = tf.assign(spike, tf.constant(True))
20         updater.eval()
21     else:
22         tf.assign(spike, False).eval()
23     print("Spike", spike.eval())
24
25 # Check to see if there some uninitialized variables:
26 print(sess.run(tf.report_uninitialized_variables()))
27
28 sess.close()
```



Python ▼

Spike	False
Spike	True
Spike	False
Spike	False
Spike	True
Spike	False
Spike	True

... try it in class

TensorFlow

- Examples:

```
1  ## Saving Variables in TensorFlow
2  import tensorflow as tf
3  sess = tf.InteractiveSession()
4
5  # Create a boolean vector called `spike` to locate a sudden spike in data.
6  # Since all variables must be initialized, initialize the variable by calling
7  # `run()` on its `initializer`.
8  vector = [3., -2., 8., -4., 0.2, 2.3, 7.5, 14.8]
9  spikes = tf.Variable([False] * len(vector), name='spikes')
10
11 # Initiallizing our variable in one of two ways (choose one):
12 # sess.run(spikes.initializer)
13 spikes.initializer.run()
14
15 # The saver op will enable saving and restoring
16 saver = tf.train.Saver()
17
18 # Loop through the data and update the spike variable when there is a significant
19 # increase
20 for i in range(1, len(vector)):
21     if vector[i] - vector[i-1] > 5:
22         spikes_val = spikes.eval()
23         spikes_val[i] = True
24         updater = tf.assign(spikes, spikes_val)
25         updater.eval()
26
27 save_path = saver.save(sess, "./spikes.ckpt")
28 print("spikes data saved in file: %s" % save_path)
29
30 sess.close()
```

spikes data saved in file: ./spikes.ckpt

- checkpoint
- spikes.ckpt.data-00000-of-00001
- spikes.ckpt.index
- spikes.ckpt.meta

... try it in class

TensorFlow

- Examples:

```
1  ## Loading Variables in TensorFlow
2  import tensorflow as tf
3  sess = tf.InteractiveSession()
4
5  # Create a boolean vector called `spike` to locate a sudden spike in data.
6  # Since all variables must be initialized, initialize the variable by calling
7  # `run()` on its `initializer`.
8  spikes = tf.Variable([False]*8, name='spikes')
9  saver = tf.train.Saver()
10
11  saver.restore(sess, "./spikes.ckpt")
12  print(spikes.eval())
13
14  sess.close()
```

```
INFO:tensorflow:Restoring parameters from ./spikes.ckpt
[False False  True False False False  True  True]
```

TensorFlow

- Examples:

```
1  ## Log example:
2  import tensorflow as tf
3
4  # Let's create a simple matrix:
5  matrix = tf.constant([[3., 4.]])
6
7  # Now negate it:
8  negMatrix = tf.negative(matrix)
9
10 # Let's see where each operation is mapped to:
11 with tf.Session(config=tf.ConfigProto(log_device_placement=True)) as sess:
12     result = sess.run(negMatrix)
13
14 # Print results on screen:
15 print(result)
16 print(matrix.shape)
17
18 # To access each member inside a tensor do:
19 print(matrix.shape[0])
```


```
2018-05-28 14:51:29.346491: I tensorflow/core/common_runtime/direct_session.cc:297] Device mapping:
2018-05-28 14:51:29.347240: I tensorflow/core/common_runtime/placer.cc:874] Neg: (Neg)/job:localhost/
replica:0/task:0/device:CPU:0
[[-3. -4.]] 2018-05-28 14:51:29.347253: I tensorflow/core/common_runtime/placer.cc:874] Const: (Const)
/job:localhost/replica:0/task:0/device:CPU:0
```

```
(1, 2)
1
```

... try it in class

NumPy arrays recap

- Difference between NumPy arrays vs Python Lists
 - NumPy array:
 - A **NumPy array** is a Python object **build around a C array**
 - This means that it has a pointer to a **contiguous data buffer of values**
 - Python Lists:
 - A **Python list** has a pointer to a **contiguous buffer of pointers**
 - **All of them point to different Python objects**, which in turn has references to its data (in this case, integers)
 - Conclusion:
 - NumPy is much more efficient than Python, in the **cost of storage** and in **speed of access**



NumPy arrays recap

- NumPy arrays

Example:

... try it in class

```
Python
In [23]: a = np.array([[12, 34, 41], [54, 62, 18], [72, 84, 96]], np.int16)

In [24]: a
Out[24]:
array([[12, 34, 41],
       [54, 62, 18],
       [72, 84, 96]], dtype=int16)

In [25]: a.size
Out[25]: 9

In [26]: a.shape
Out[26]: (3, 3)

In [27]: type(a)
Out[27]: numpy.ndarray

In [28]: a.dtype
Out[28]: dtype('int16')

In [29]: a[2,2] # this is how we index a particular elemnt in the array (#9)
Out[29]: 96

In [30]: b = a[0,:]

In [31]: b
Out[31]: array([12, 34, 41], dtype=int16)

In [32]: b.shape
Out[32]: (3,)

In [33]: b[2] = 88 # this is how we reassign another value to a member in the array

In [34]: a[2,2] = 99 # the change above also affects the original array

In [35]: a
Out[35]:
array([[12, 34, 88],
       [54, 62, 18],
       [72, 84, 99]], dtype=int16)


In [36]: b
Out[36]: array([12, 34, 88], dtype=int16)
```

TensorFlow

- Let's discuss the basics of **computation graphs** without the context of Tensor-Flow
- This includes:
 - defining **nodes**
 - defining **edges**
 - **dependencies**
 - examples to illustrate **key principles**

TensorFlow


- Graph basics:
 - At **the core** of every TensorFlow program is the **computation graph**
 - It is a specific type of directed graph that is used for **defining computational structure**
 - In TensorFlow it is, a **series of functions chained together**, each passing its output to zero, one, or more functions further along in the chain



$$f(1, 2) = 1 + 2 = 3$$

TensorFlow


- Graph basics:
 - **Nodes**: typically drawn as circles, ovals, or boxes, represent some sort of computation or action being done on or with data in the graph's context. In the example below, the operation “add” is the sole node.



$$f(1, 2) = 1 + 2 = 3$$


TensorFlow

- Graph basics:
 - **Edges**: are the actual values that get passed to and from *Operations*, and are typically drawn as **arrows**
 - In the “add” example, the inputs 1 and 2 are both edges leading into the node, while the output 3 is an edge leading out of the node
 - We can think of edges as the link between different *Operations* as **they carry information** from one node to the next



$$f(1, 2) = 1 + 2 = 3$$

TensorFlow



TensorFlow

We can decompose this graphical representation as a series of equations like this:

$$a = \text{input}_1; b = \text{input}_2$$

$$c = a \cdot b; d = a + b$$

$$e = c + d$$


to solve e for $a = 5$ and $b = 3$,

$$a = 5; b = 3$$


$$e = (a \cdot b) + (a + b)$$

$$e = (5 \cdot 3) + (5 + 3)$$

$$e = 15 + 8 = 23$$




TensorFlow



TensorFlow


- Graph basics:
 - **Dependencies**: there are certain types of connections between nodes that aren't allowed, the most common of which is one that creates an unresolved *circular dependency*



TensorFlow

- Graph basics:
 - Dependencies: ...


let's take a look at what happens if the multiplication node c is unable to finish its computation (for whatever reason):



TensorFlow

- Graph basics:
 - Dependencies: ...


What happens if one of the inputs fails to pass its data on to the next functions in the graph?



TensorFlow

- Graph basics:
 - **Dependencies:** ...


Let's see what happens if we redirect the output of a graph back into an earlier portion of it:



TensorFlow

- Graph basics:
 - Because of this, truly circular dependencies can't be expressed in TensorFlow, which is not a bad thing.
 - Dependencies:** ...


Let's provide an initial state to the value feeding into either *b* or *e*. Let's give the graph a kick-start by giving the output of *e* an initial value of 1:



TensorFlow


- Graph basics:
 - **Dependencies:** ...

By unrolling our graph like this, we can simulate useful cyclical dependencies while maintaining a deterministic computation.



TensorFlow

- Building our first graph in TensorFlow:



```
simple_graph.py
1 ## Building our first TensorFlow graph:
2
3 # First we need to import TensorFlow:
4 import tensorflow as tf
5
6 # Let's define our input nodes:
7 a = tf.constant(5, name="input_a")
8 b = tf.constant(3, name="input_b")
9
10 # Defining the next two nodes in our graph:
11 c = tf.multiply(a,b, name="mul_c")
12 d = tf.add(a,b, name="add_d")
13
14 # This last line defines the final node in our graph:
15 e = tf.add(c,d, name="add_e")
```

TensorFlow

- Building our first graph in TensorFlow:

```
Python /Users/alex

In [6]: whos
Variable  Type      Data/Info
-----
a         Tensor    Tensor("input_a_1:0", shape=(), dtype=int32)
b         Tensor    Tensor("input_b_1:0", shape=(), dtype=int32)
c         Tensor    Tensor("mul_c:0", shape=(), dtype=int32)
d         Tensor    Tensor("add_d:0", shape=(), dtype=int32)
e         Tensor    Tensor("add_e:0", shape=(), dtype=int32)
tf        module    <module 'tensorflow' from<...>tensorflow/__init__.pyc'>
```

```
simple_graph.py

1  ## Building our first TensorFlow graph:
2
3  # First we need to import TensorFlow:
4  import tensorflow as tf
5
6  # Let's define our input nodes:
7  a = tf.constant(5, name="input_a")
8  b = tf.constant(3, name="input_b")
9
10 # Defining the next two nodes in our graph:
11 c = tf.multiply(a,b, name="mul_c")
12 d = tf.add(a,b, name="add_d")
13
14 # This last line defines the final node in our graph:
15 e = tf.add(c,d, name="add_e")
```

To run we have to add the two extra lines and run them in the shell:

```
In [7]: sess = tf.Session()

In [8]: sess.run(e)
Out[8]: 23
```

TensorFlow

- Let's construct the actual graph using TensorBoard:
 - First, we need to make sure we have generated summary data in a log directory by creating a summary writer:

```
21 # To create the graph:  
22 sess.graph.as_graph_def()  
23 file_writer = tf.summary.FileWriter('./', sess.graph)
```

sess.graph.as_graph_def()
contains the graph definition that
enables the Graph Visualizer

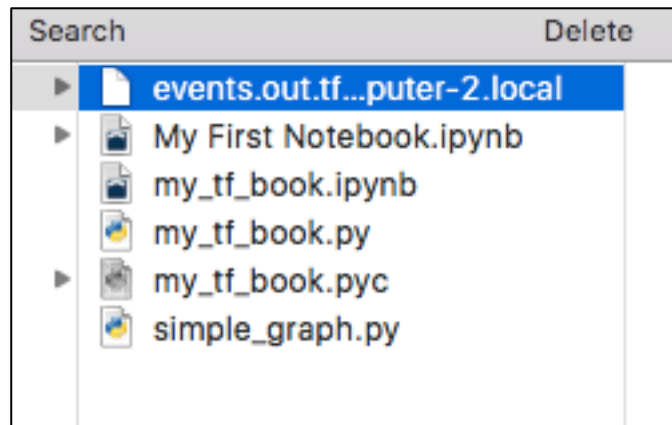
Docstring:
Writes `Summary` protocol
buffers to event files.

'./path/to/logs'

The diagram illustrates the code for creating a summary writer in TensorFlow. It features a code block with three lines. The first line is a comment. The second line calls `sess.graph.as_graph_def()`. The third line creates a `tf.summary.FileWriter` object, passing `'./'` and `sess.graph` as arguments. Three blue arrows point from explanatory text to the code: one from the top right to `sess.graph.as_graph_def()`, one from the bottom left to `tf.summary.FileWriter`, and one from the bottom right to the path `'./'`.

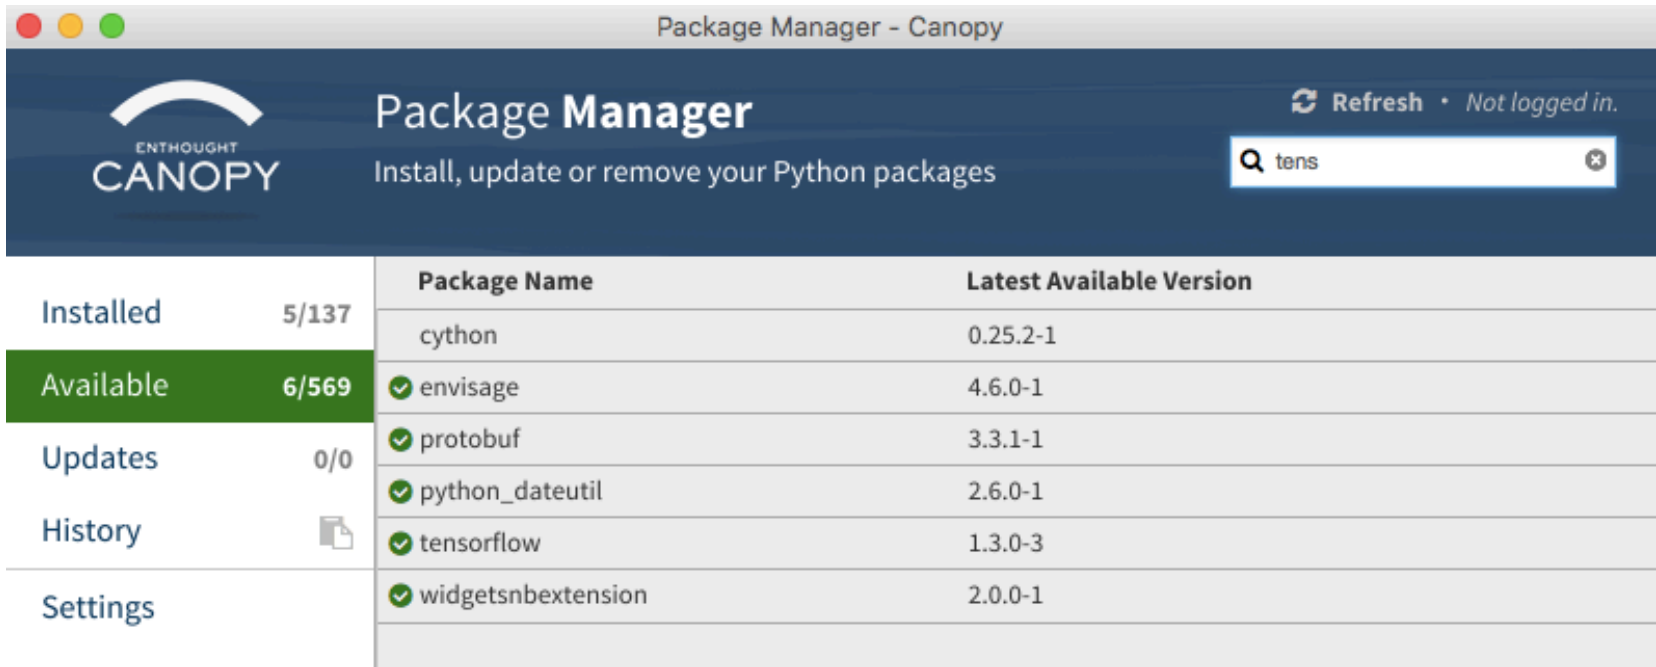
TensorFlow

- Let's construct the actual graph using TensorBoard:
- Once we run the previous code, a file with the session is generated in our current folder:



TensorFlow

- Let's construct the actual graph using TensorBoard:
 - Before we continue we need to check if we have TensorBoard installed in our system:

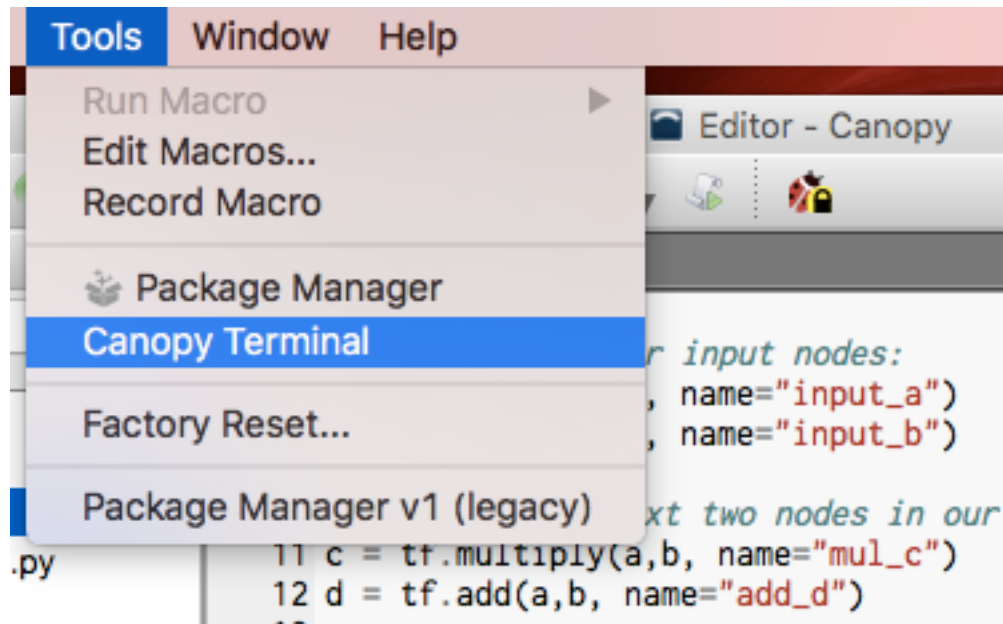


The screenshot shows the Canopy Package Manager window. The title bar reads "Package Manager - Canopy". The header area includes the Canopy logo, the text "Package Manager", and a subtitle "Install, update or remove your Python packages". There is a "Refresh" button and a status indicator "Not logged in.". A search bar contains the text "tens".

		Package Name	Latest Available Version
Installed	5/137	cython	0.25.2-1
Available	6/569	✓ envisage	4.6.0-1
		✓ protobuf	3.3.1-1
		✓ python_dateutil	2.6.0-1
		✓ tensorflow	1.3.0-3
		✓ widgetsnbextension	2.0.0-1

TensorFlow

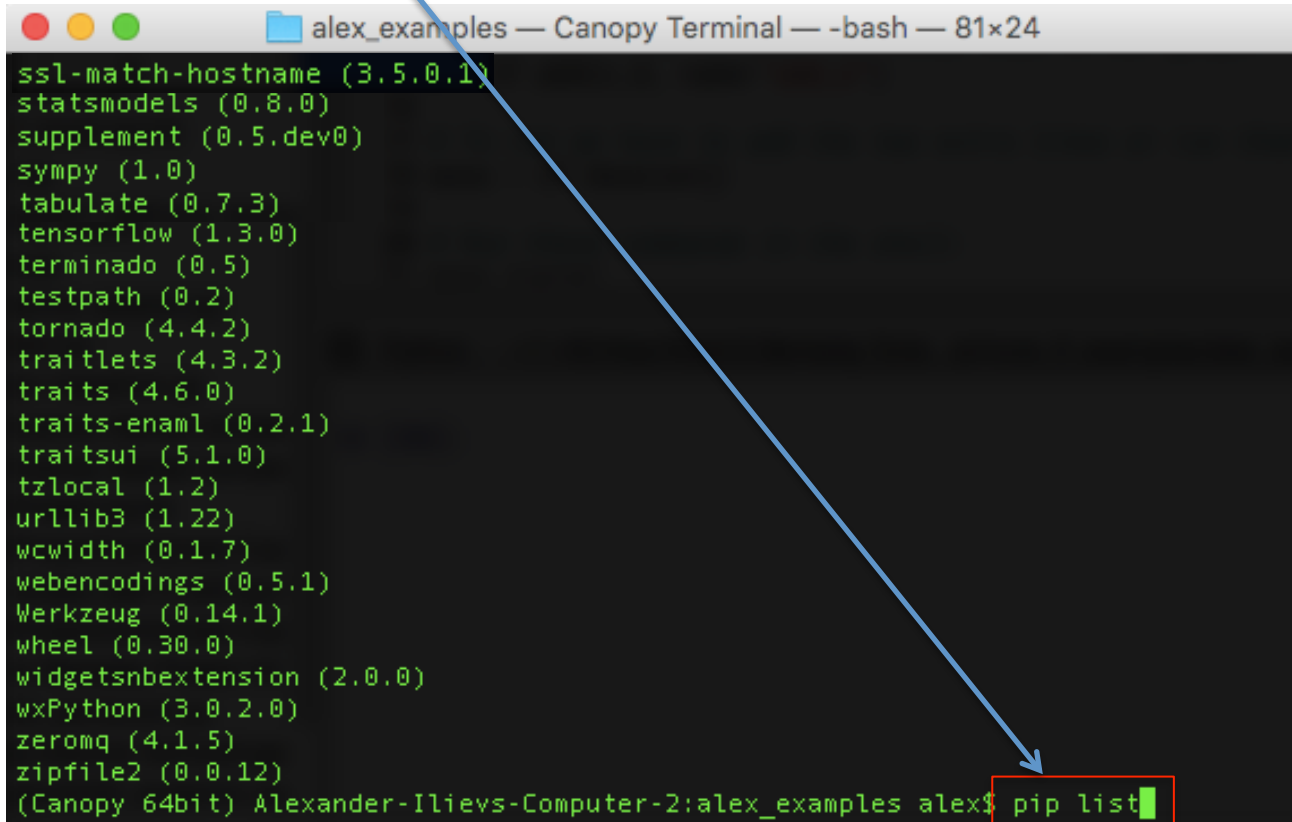
- Let's construct the actual graph using TensorBoard:
- Canopy **does not provide TensorBoard** in its repository, therefore we need to install it via the Canopy Terminal:



TensorFlow

- Let's construct the actual graph using TensorBoard:
- We then type: **pip list** to check for installed packages:

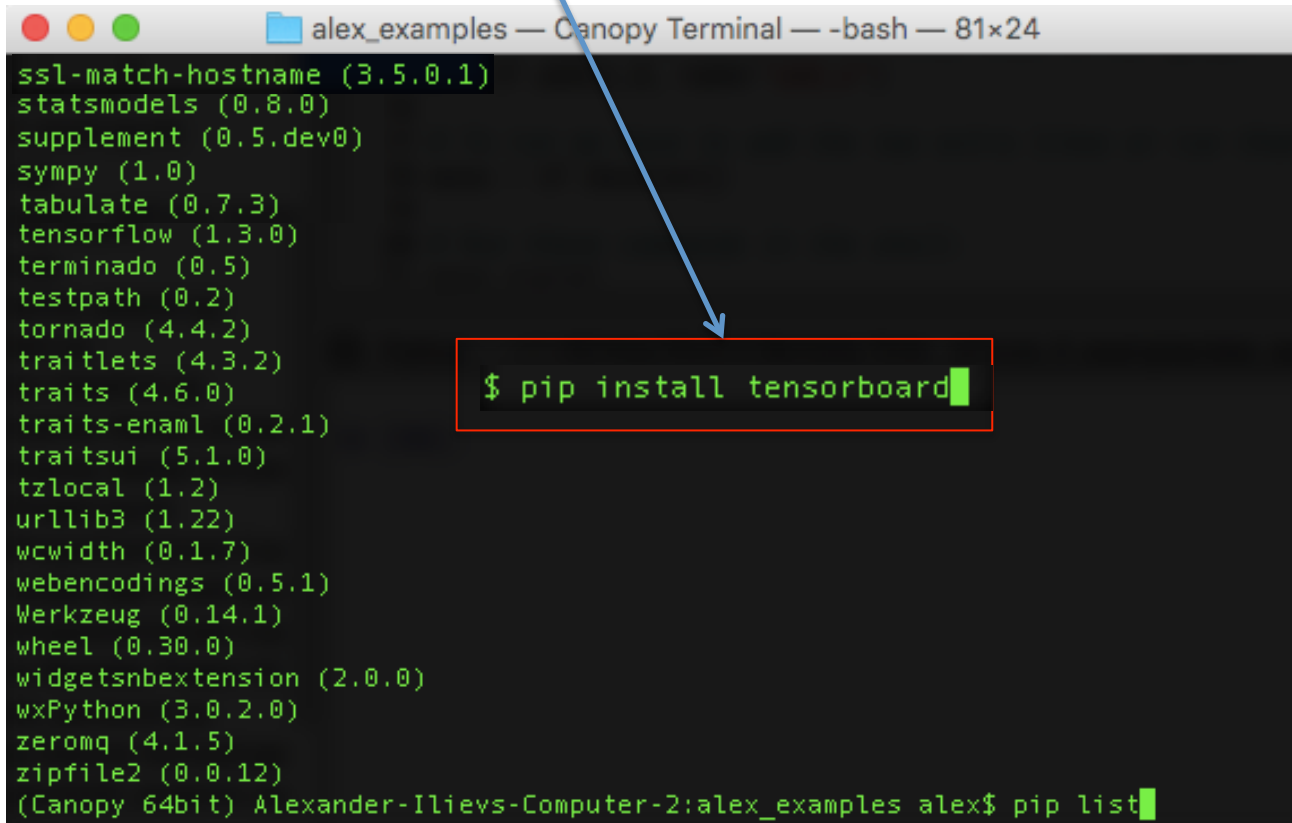
TensorBoard is not
found in the list



```
alex_examples — Canopy Terminal — -bash — 81x24
ssl-match-hostname (3.5.0.1)
statsmodels (0.8.0)
supplement (0.5.dev0)
sympy (1.0)
tabulate (0.7.3)
tensorflow (1.3.0)
terminado (0.5)
testpath (0.2)
tornado (4.4.2)
traitlets (4.3.2)
traits (4.6.0)
traits-enaml (0.2.1)
traitsui (5.1.0)
tzlocal (1.2)
urllib3 (1.22)
wcwidth (0.1.7)
webencodings (0.5.1)
Werkzeug (0.14.1)
wheel (0.30.0)
widgetsnbextension (2.0.0)
wxPython (3.0.2.0)
zeromq (4.1.5)
zipfile2 (0.0.12)
(Canopy 64bit) Alexander-Iliev-Computer-2:alex_examples alex$ pip list
```

TensorFlow

- Let's construct the actual graph using TensorBoard:
- We then type: **pip install tensorboard** in the terminal

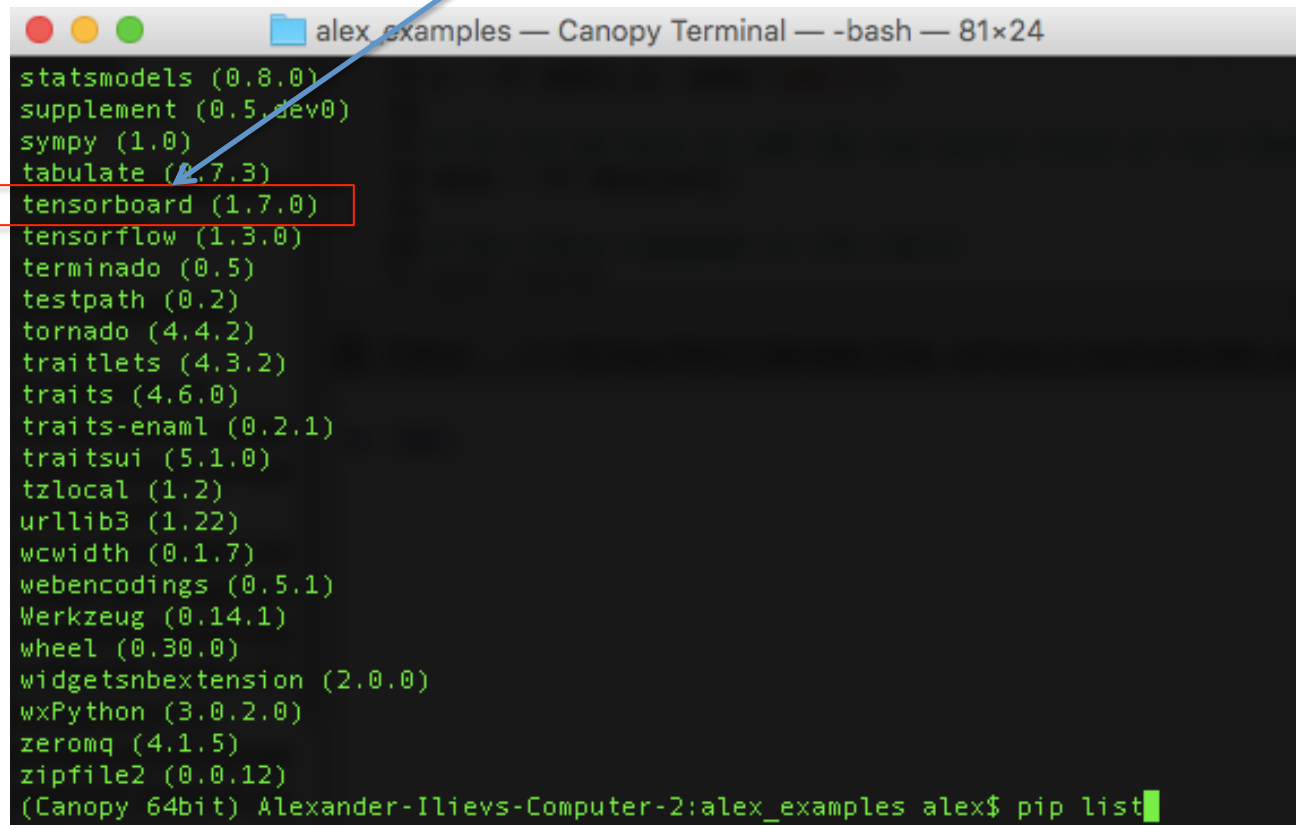


A screenshot of a terminal window titled "alex_examples — Canopy Terminal — -bash — 81x24". The terminal displays a list of installed packages with their versions: ssl-match-hostname (3.5.0.1), statsmodels (0.8.0), supplement (0.5.dev0), sympy (1.0), tabulate (0.7.3), tensorflow (1.3.0), terminado (0.5), testpath (0.2), tornado (4.4.2), traitlets (4.3.2), traits (4.6.0), traits-enaml (0.2.1), traitsui (5.1.0), tzlocal (1.2), urllib3 (1.22), wcwidth (0.1.7), webencodings (0.5.1), Werkzeug (0.14.1), wheel (0.30.0), widgetsnbextension (2.0.0), wxPython (3.0.2.0), zeromq (4.1.5), and zipfile2 (0.0.12). At the bottom, the prompt "(Canopy 64bit) Alexander-Iliev's-Computer-2:alex_examples alex\$ pip list" is shown. A blue arrow points from the text "pip install tensorboard" in the list above to a red-bordered box containing the command "\$ pip install tensorboard" in the terminal.

```
alex_examples — Canopy Terminal — -bash — 81x24
ssl-match-hostname (3.5.0.1)
statsmodels (0.8.0)
supplement (0.5.dev0)
sympy (1.0)
tabulate (0.7.3)
tensorflow (1.3.0)
terminado (0.5)
testpath (0.2)
tornado (4.4.2)
traitlets (4.3.2)
traits (4.6.0)
traits-enaml (0.2.1)
traitsui (5.1.0)
tzlocal (1.2)
urllib3 (1.22)
wcwidth (0.1.7)
webencodings (0.5.1)
Werkzeug (0.14.1)
wheel (0.30.0)
widgetsnbextension (2.0.0)
wxPython (3.0.2.0)
zeromq (4.1.5)
zipfile2 (0.0.12)
(Canopy 64bit) Alexander-Iliev's-Computer-2:alex_examples alex$ pip list
```

TensorFlow

- Let's construct the actual graph using TensorBoard:
- We check again: **tensorboard** is now installed



```
statsmodels (0.8.0)
supplement (0.5.dev0)
sympy (1.0)
tabulate (2.7.3)
tensorboard (1.7.0)
tensorflow (1.3.0)
terminado (0.5)
testpath (0.2)
tornado (4.4.2)
traitlets (4.3.2)
traits (4.6.0)
traits-enaml (0.2.1)
traitsui (5.1.0)
tzlocal (1.2)
urllib3 (1.22)
wcwidth (0.1.7)
webencodings (0.5.1)
Werkzeug (0.14.1)
wheel (0.30.0)
widgetsnbextension (2.0.0)
wxPython (3.0.2.0)
zeromq (4.1.5)
zipfile2 (0.0.12)
(Canopy 64bit) Alexander-Iliev-Computer-2:alex_examples alex$ pip list
```

TensorFlow

- Let's construct the actual graph using TensorBoard:
 - WARNING:** We check for compatibility: **tensorflow** + **tensorboard**

```
tensorboard 1.8.0
tensorflow 1.6.0
```

When installing tensorboard you might get this message:

```
tensorflow 1.6.0 has requirement tensorboard<1.7.0,>=1.6.0, but you'll have tensorboard 1.8.0 which is incompatible.
```

When running tensorboard you might get this message:

```
(Canopy 64bit) Alexandomputer2:alex_examples alex$ tensorboard --logdir ./ --event_file events.out.tfevents.1527382420.Alexandomputer2
2018-05-26 18:00:37.024647: I tensorflow/core/platform/cpu_feature_guard.cc:140]
Your CPU supports instructions that this TensorFlow binary was not compiled to
use: SSE4.1
```

So you have to install a compatible version:

```
alex$ pip install tensorboard==1.6.0
```

Make sure they match:

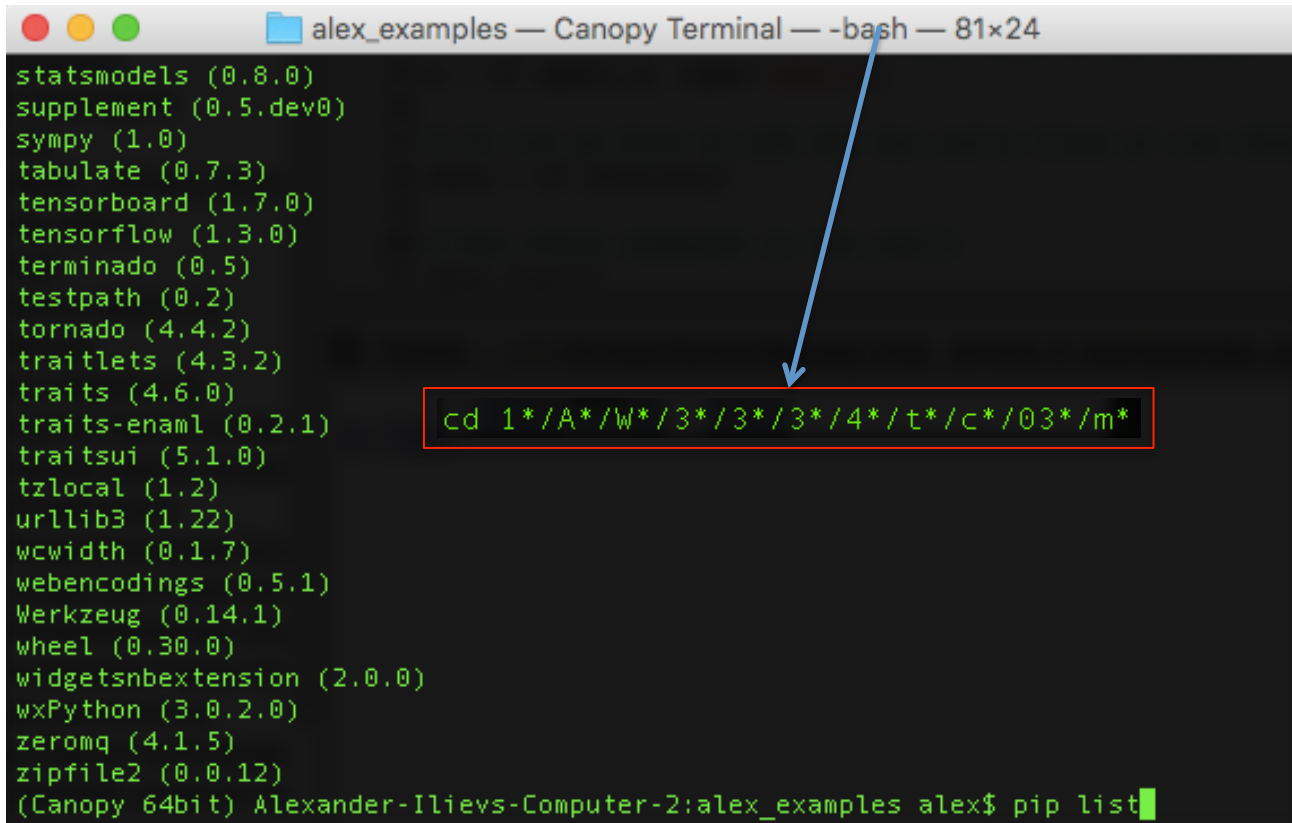
```
tensorboard 1.6.0
tensorflow 1.6.0
```

When running tensorboard it runs without warnings:

```
(Canopy 64bit) Alexandomputer2:alex_examples alex$ tensorboard --logdir ./ --event_file events.out.tfevents.1527382420.Alexandomputer2
TensorBoard 1.6.0 at http://Alexandomputer2:6006 (Press CTRL+C to quit)
```

TensorFlow

- Let's construct the actual graph using TensorBoard:
- We must be operating in the correct **directory** before we go on:

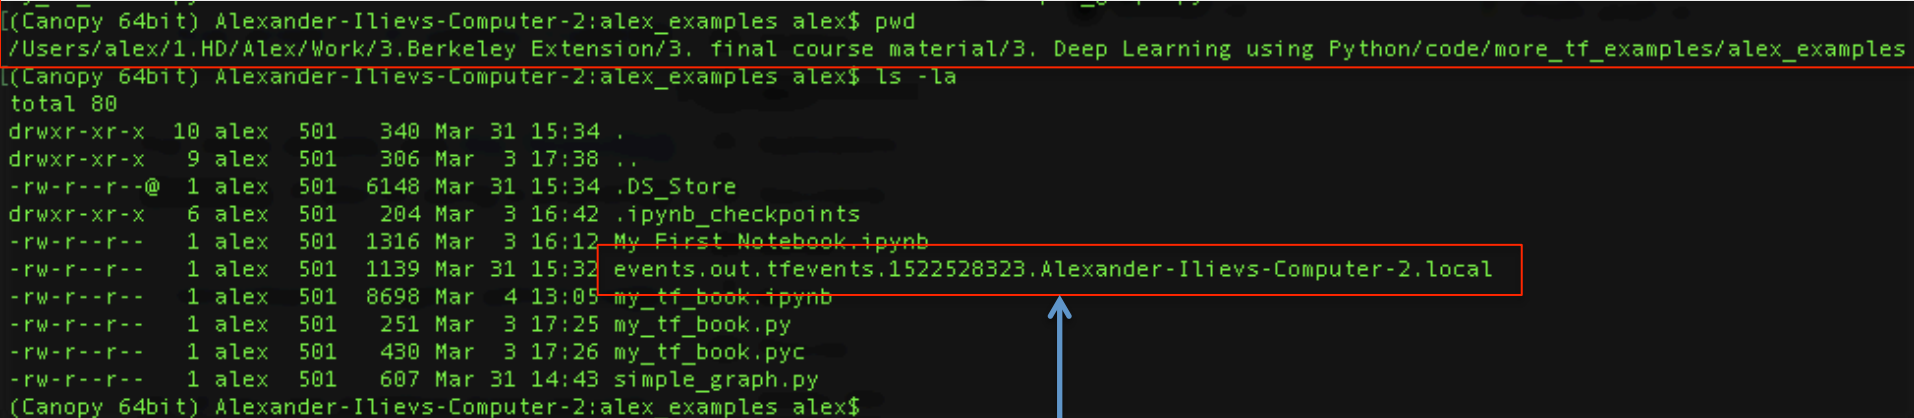


The screenshot shows a Canopy Terminal window with the title bar "alex_examples — Canopy Terminal — -bash — 81x24". The terminal displays a list of installed packages and their versions. A blue arrow points from the word "directory" in the text above to a red-bordered box containing the command `cd 1*/A*/W*/3*/3*/3*/4*/t*/c*/03*/m*`. At the bottom of the terminal, the command `(Canopy 64bit) Alexander-Iliev-Computer-2:alex_examples alex$ pip list` is entered.

```
statsmodels (0.8.0)
supplement (0.5.dev0)
sympy (1.0)
tabulate (0.7.3)
tensorboard (1.7.0)
tensorflow (1.3.0)
terminado (0.5)
testpath (0.2)
tornado (4.4.2)
traitlets (4.3.2)
traits (4.6.0)
traits-enaml (0.2.1)
traitsui (5.1.0)
tzlocal (1.2)
urllib3 (1.22)
wcwidth (0.1.7)
webencodings (0.5.1)
Werkzeug (0.14.1)
wheel (0.30.0)
widgetsnbextension (2.0.0)
wxPython (3.0.2.0)
zeromq (4.1.5)
zipfile2 (0.0.12)
(Canopy 64bit) Alexander-Iliev-Computer-2:alex_examples alex$ pip list
```

TensorFlow

- Let's construct the actual graph using TensorBoard:
 - We must be operating in the correct **directory** before we go on:



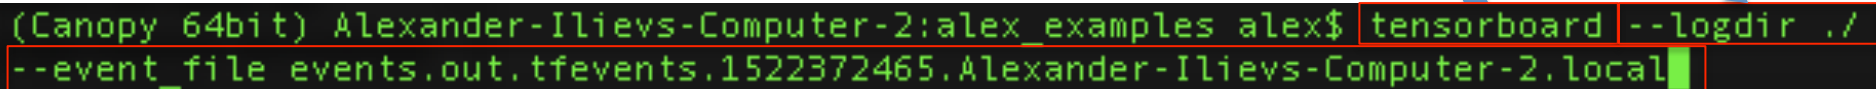
A terminal window showing the current directory and its contents. A blue arrow points from the word 'directory' in the text above to the terminal output. A red box highlights the file 'events.out.tfevents.1522528323.Alexander-Ilievs-Computer-2.local', and a blue arrow points from the text below to this file.

```
(Canopy 64bit) Alexander-Ilievs-Computer-2:alex_examples alex$ pwd
/Users/alex/1.HD/Alex/Work/3.Berkeley Extension/3. final course material/3. Deep Learning using Python/code/more_tf_examples/alex_examples
(Canopy 64bit) Alexander-Ilievs-Computer-2:alex_examples alex$ ls -la
total 80
drwxr-xr-x  10 alex  501   340 Mar 31 15:34 .
drwxr-xr-x   9 alex  501   306 Mar  3 17:38 ..
-rw-r--r--@  1 alex  501  6148 Mar 31 15:34 .DS_Store
drwxr-xr-x   6 alex  501   204 Mar  3 16:42 .ipynb_checkpoints
-rw-r--r--   1 alex  501  1316 Mar  3 16:12 My First Notebook.ipynb
-rw-r--r--   1 alex  501  1139 Mar 31 15:32 events.out.tfevents.1522528323.Alexander-Ilievs-Computer-2.local
-rw-r--r--   1 alex  501  8698 Mar  4 13:05 my_tf_book.ipynb
-rw-r--r--   1 alex  501   251 Mar  3 17:25 my_tf_book.py
-rw-r--r--   1 alex  501   430 Mar  3 17:26 my_tf_book.pyc
-rw-r--r--   1 alex  501   607 Mar 31 14:43 simple_graph.py
(Canopy 64bit) Alexander-Ilievs-Computer-2:alex_examples alex$
```

- Then we make sure the **graph summary file** was created

TensorFlow

- Let's construct the actual graph using TensorBoard:
 - Once we have the event file(s), **we run TensorBoard** while providing the **log directory**:

A terminal window showing a command to run TensorBoard. The command is `tensorboard --logdir ./ --event_file events.out.tfevents.1522372465.Alexander-Ilievs-Computer-2.local`. Two blue arrows point from the text above to the command: one points to `tensorboard` and the other points to `--logdir ./`.

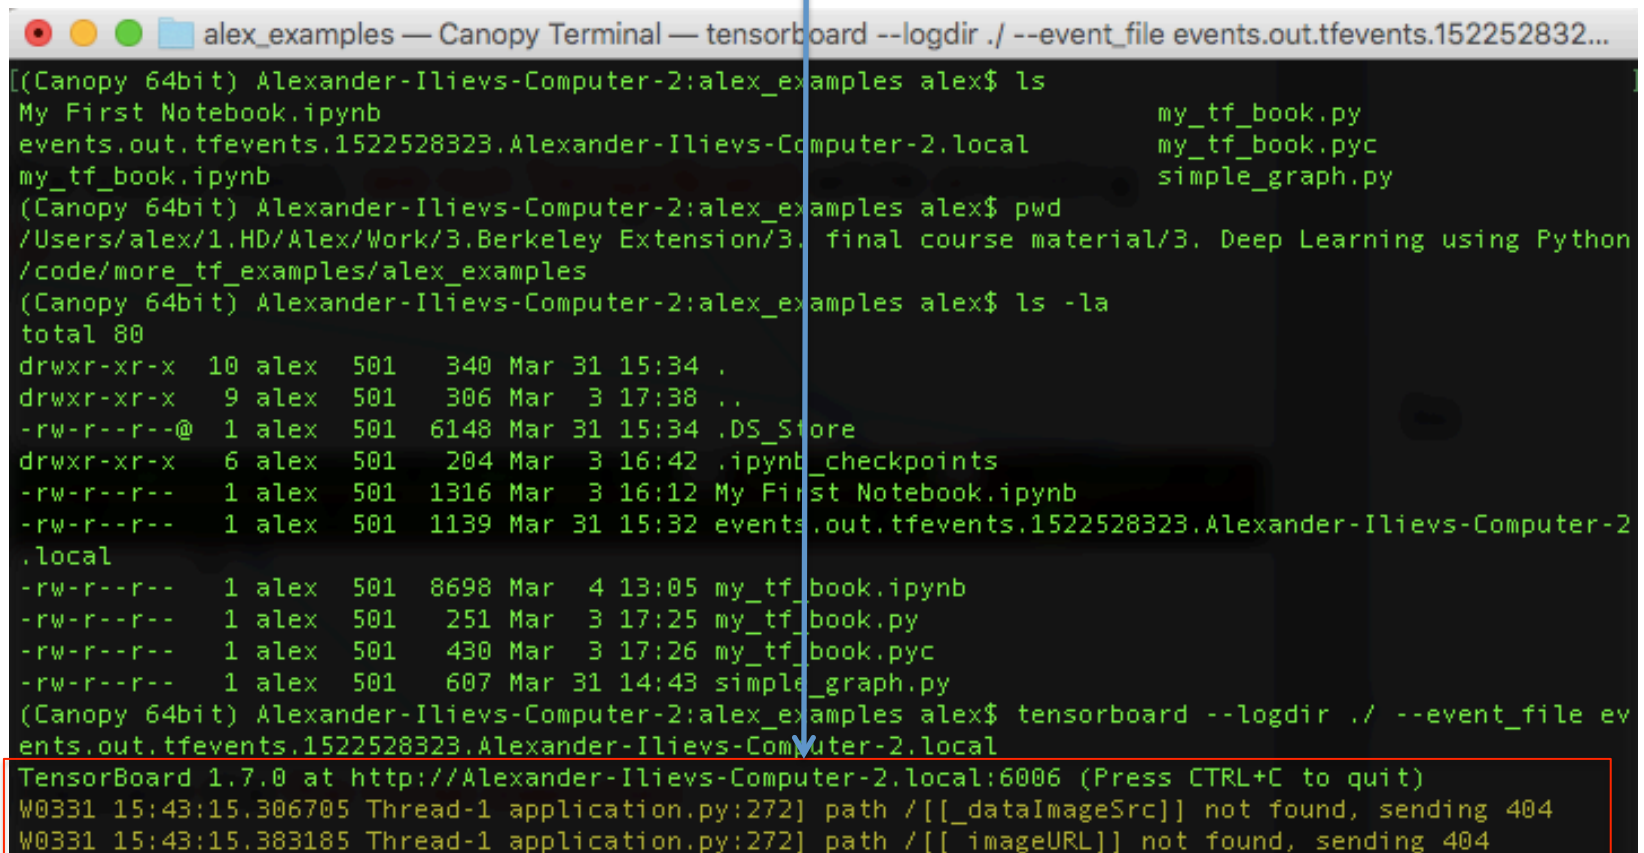
```
(Canopy 64bit) Alexander-Ilievs-Computer-2:alex_examples alex$ tensorboard --logdir ./  
--event_file events.out.tfevents.1522372465.Alexander-Ilievs-Computer-2.local
```

- and specifically request the **file to be executed**
- or if you are in the same directory where the file resides simply run:

```
tensorboard --logdir ./
```


TensorFlow


- Let's construct the actual graph using TensorBoard:
 - And the graph is already **running** in the background:



```
alex_examples — Canopy Terminal — tensorboard --logdir ./ --event_file events.out.tfevents.1522528323...  
[(Canopy 64bit) Alexander-Ilievs-Computer-2:alex_examples alex$ ls  
My First Notebook.ipynb  
events.out.tfevents.1522528323.Alexander-Ilievs-Computer-2.local  
my_tf_book.ipynb  
(Canopy 64bit) Alexander-Ilievs-Computer-2:alex_examples alex$ pwd  
/Users/alex/1.HD/Alex/Work/3.Berkeley Extension/3. final course material/3. Deep Learning using Python  
/code/more_tf_examples/alex_examples  
(Canopy 64bit) Alexander-Ilievs-Computer-2:alex_examples alex$ ls -la  
total 80  
drwxr-xr-x  10 alex  501   340 Mar 31 15:34 .  
drwxr-xr-x   9 alex  501   306 Mar  3 17:38 ..  
-rw-r--r--@  1 alex  501  6148 Mar 31 15:34 .DS_Store  
drwxr-xr-x   6 alex  501   204 Mar  3 16:42 .ipynb_checkpoints  
-rw-r--r--   1 alex  501  1316 Mar  3 16:12 My First Notebook.ipynb  
-rw-r--r--   1 alex  501  1139 Mar 31 15:32 events.out.tfevents.1522528323.Alexander-Ilievs-Computer-2  
.local  
-rw-r--r--   1 alex  501  8698 Mar  4 13:05 my_tf_book.ipynb  
-rw-r--r--   1 alex  501   251 Mar  3 17:25 my_tf_book.py  
-rw-r--r--   1 alex  501   430 Mar  3 17:26 my_tf_book.pyc  
-rw-r--r--   1 alex  501   607 Mar 31 14:43 simple_graph.py  
(Canopy 64bit) Alexander-Ilievs-Computer-2:alex_examples alex$ tensorboard --logdir ./ --event_file ev  
ents.out.tfevents.1522528323.Alexander-Ilievs-Computer-2.local  
TensorBoard 1.7.0 at http://Alexander-Ilievs-Computer-2.local:6006 (Press CTRL+C to quit)  
W0331 15:43:15.306705 Thread-1 application.py:272] path /[_dataImageSrc] not found, sending 404  
W0331 15:43:15.383185 Thread-1 application.py:272] path /[_imageURL] not found, sending 404
```

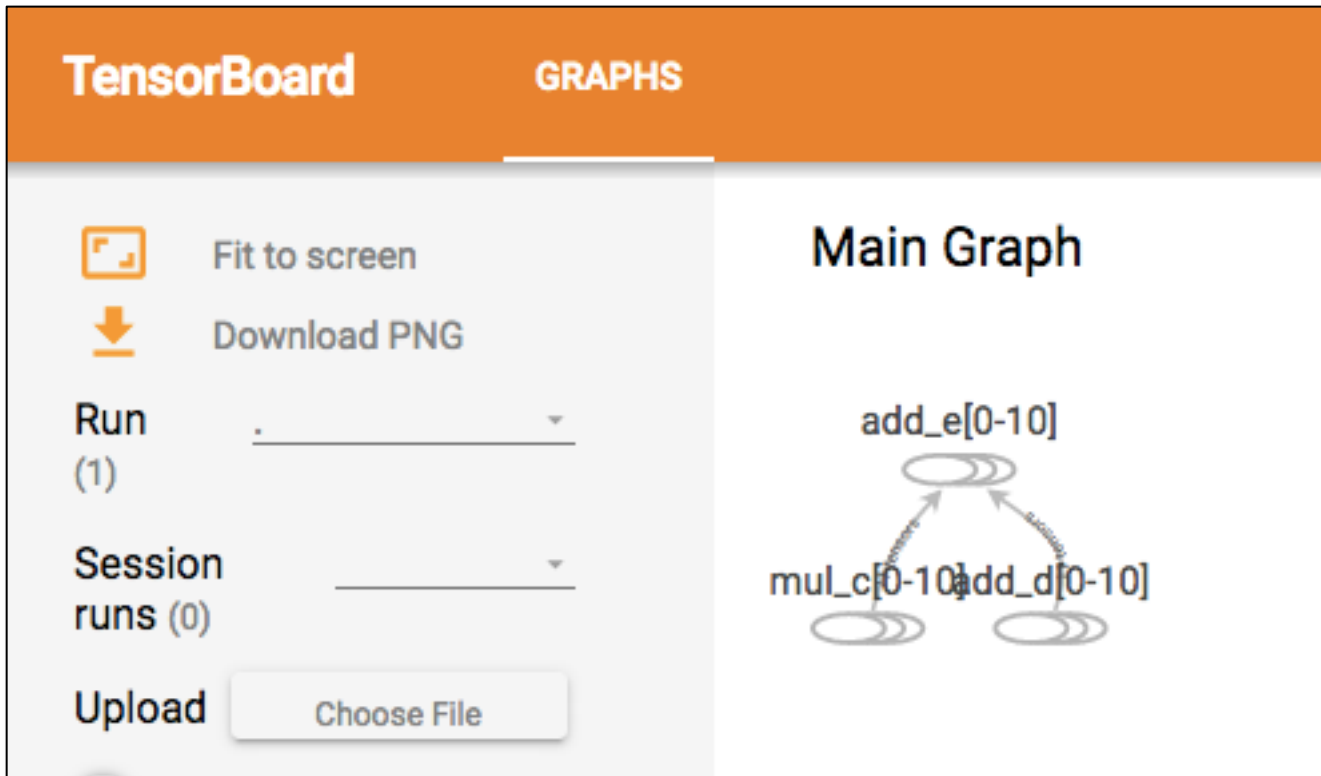
TensorFlow

- Let's construct the actual graph using TensorBoard:
- Next, we open our browser and **take a look at the graph by typing:**



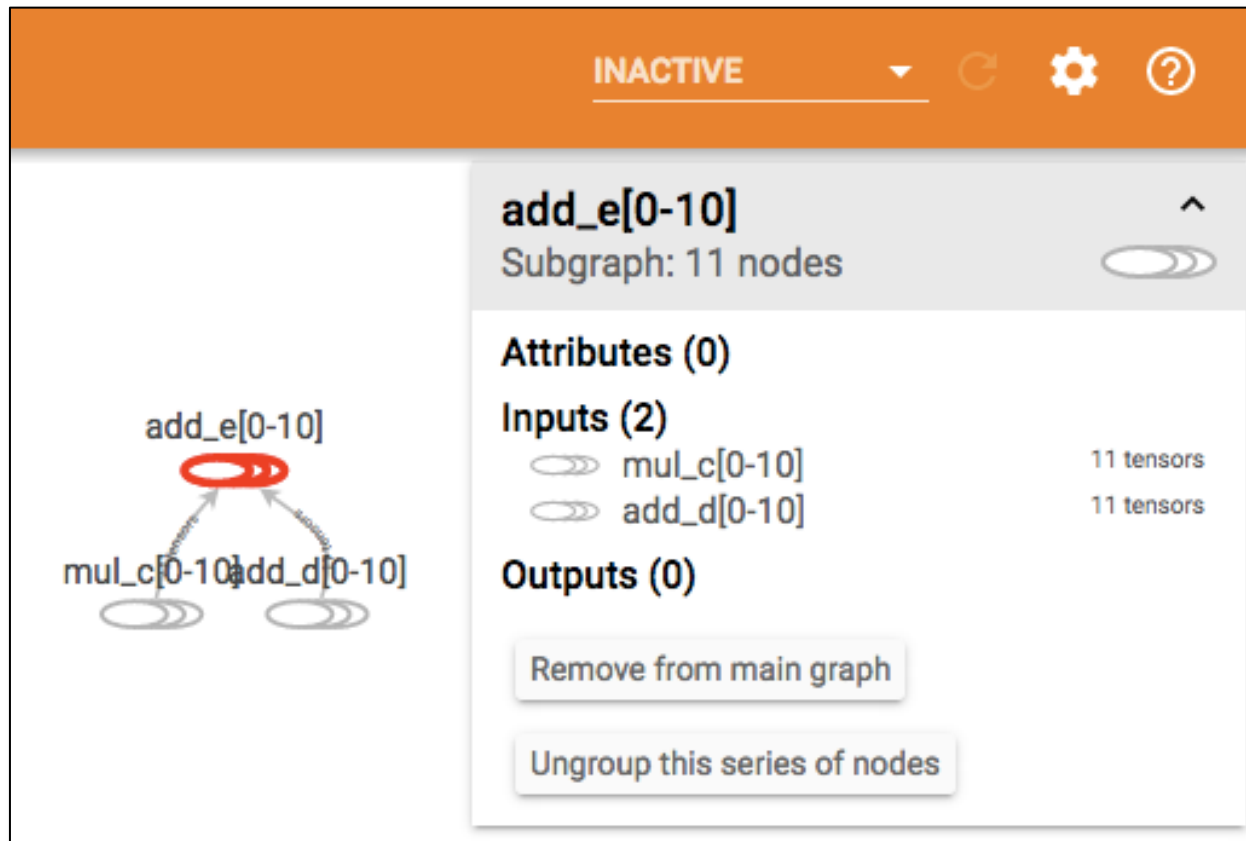
TensorFlow

- Let's construct the actual graph using TensorBoard:
- Below is a graphical representation of our first **graph**:



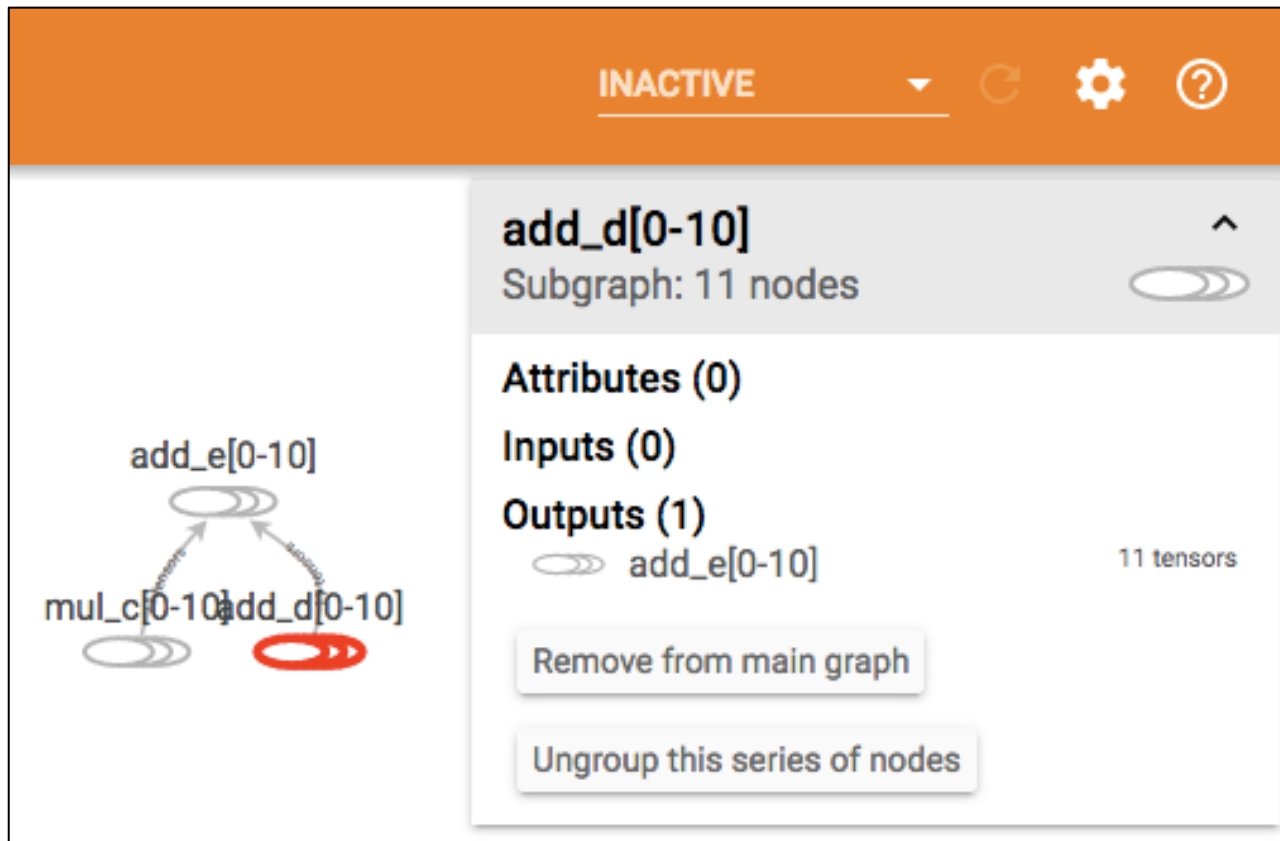
TensorFlow

- Let's construct the actual graph using TensorBoard:
- Simply click on any of the **nodes** to inspect them more closely:



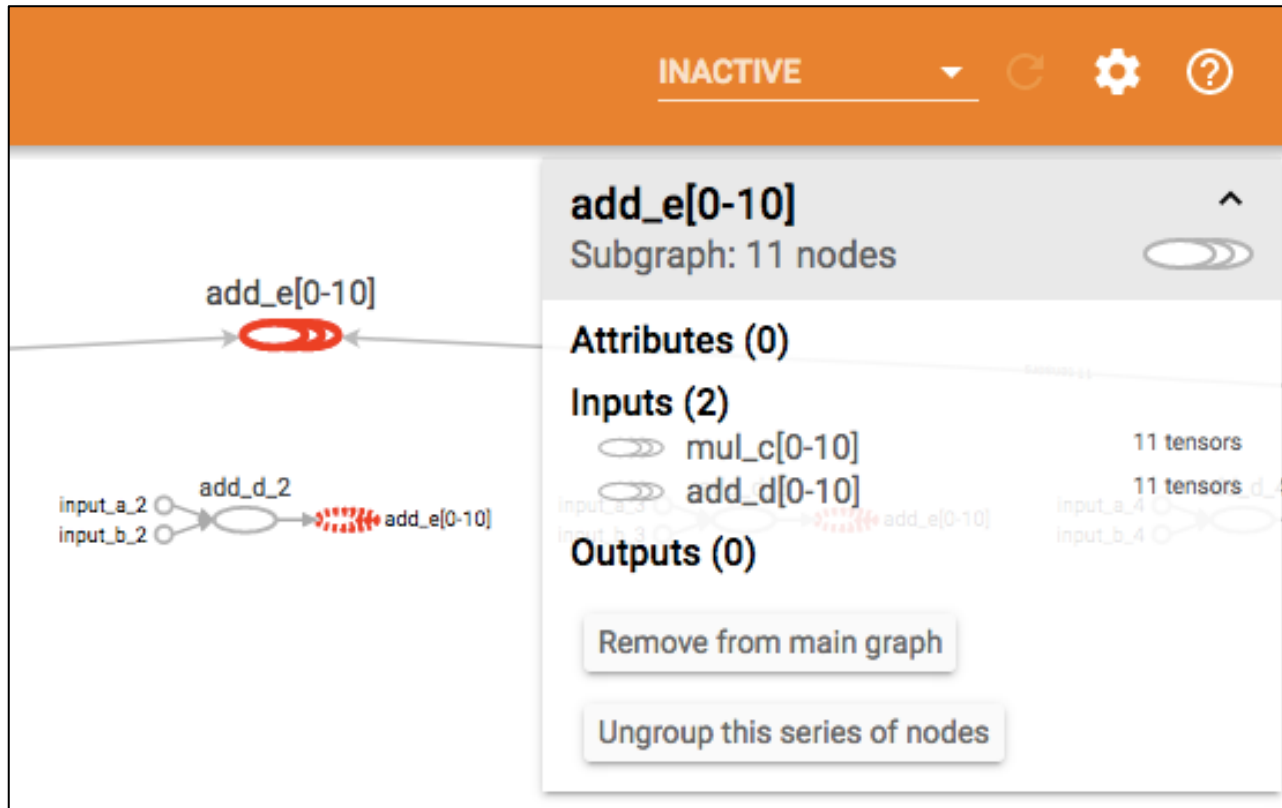
TensorFlow

- Let's construct the actual graph using TensorBoard:
- Simply click on any of the **nodes** to inspect them more closely:



TensorFlow

- Let's construct the actual graph using TensorBoard:
- Simply click on any of the **nodes** to inspect them more closely:



TensorFlow

- Let's construct the actual graph using TensorBoard:
 - Once you are done constructing our graph, we need to clean up and close the *file_writer* and *sess*:

```
25 # We clean up before we exit:  
26 file_writer.close()  
27 sess.close()
```

- In general, *Session* objects close automatically when the program terminates (or, in the interactive case, when you close/restart the Python kernel)
- However, **it's best to explicitly close out** of the *Session* to avoid any sort of weird edge case scenarios.

TensorFlow

- Instead of having two separate input nodes, we can have a **single input node** that can take in a **vector** (or **1-D tensor**) of numbers
- This graph has **several advantages** over our previous example:
 1. The client only has to send input to a **single node**, which **simplifies** using **the graph**
 2. The nodes that directly depend on the input now only have to keep track of **one dependency instead of two**
 3. We now have the option of making the graph take in vectors of any length, if we'd like. This would make the **graph more flexible**