

Machine Learning With TensorFlow

X433.7-001 (2 semester units in COMPSCI)

Instructor Alexander I. Iliev, Ph.D.

Course Content Outline

- **Machine Learning With TensorFlow®**
 - Introduction, Python - pros and cons
 - Python modules, DL packages and scientific blocks
 - Working with the shell, IPython and the editor
 - Installing the environment with core packages
 - Writing “Hello World”
- HW1 (10pts)
- **Tensorflow and TensorBoard basics**
 - Linear algebra recap
 - Data types in Numpy and Tensorflow
 - Basic operations in Tensorflow
 - Graph models and structures with Tensorboard
- **TensorFlow operations**
 - Overloaded operators
 - Using Aliases
 - Sessions, graphs, variables, placeholders
 - Name scopes
- **Data Mining and Machine Learning concepts**
 - Basic Deep Learning Models
 - Linear and Logistic Regression
 - Softmax classification
- HW2 (10pts)
- **Neural Networks**
 - Multi-layer Neuaral Network
 - Gradient descent and Backpropagation
 - Object recognition with Convolutional Neural Network (CNN)
 - Activation Functions

Data Mining

Class 3 ...

TensorFlow basics ...

TensorFlow operations

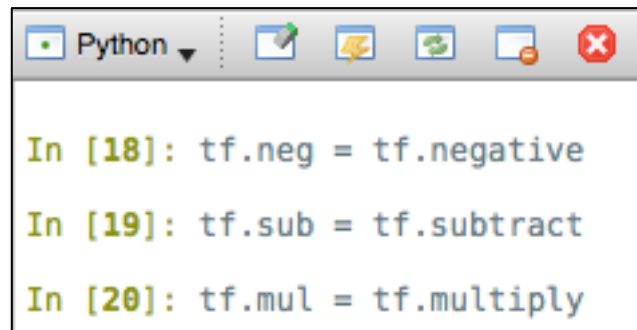
- Remember: **Tensors** are just a superset of **matrices**!
- **TensorFlow** Operations, aka Ops, **are nodes** that perform computations on or with Tensor objects
- After computation, they return zero or more tensors, which can be used by other Ops later in the graph
- **To create an Operation**, you **call its constructor** in Python
- The Python **constructor returns a handle** to the Operation's output, then it is passed on to other **Ops** or **Session.run**

TensorFlow operations

- Overloaded operators:
 - TensorFlow also **overloads** common mathematical operators to make **multiplication**, **addition**, **subtraction**, and **other operations** more concise
 - If one or more arguments to the operator is a Tensor object, a TensorFlow Operation will be called and **added to the graph**

TensorFlow operations

- Creating aliases:
 - TensorFlow allows us to create **aliases** for all common mathematical operators such as **negation**, **subtraction**, **multiplication**, or **other operations** to be more concise
 - For example, you can easily create aliases like this:



A screenshot of a Python terminal window. The title bar shows 'Python' and several icons. The terminal displays three lines of code where TensorFlow operations are aliased to shorter names:

```
In [18]: tf.neg = tf.negative
In [19]: tf.sub = tf.subtract
In [20]: tf.mul = tf.multiply
```

TensorFlow operations

- Overloaded operators:
 - Using these overloaded operators **can be great** when quickly putting together code
 - Technically, the **==** operator is overloaded as well, **but it will not return a Tensor of boolean values**. It will return **True** if the two tensors being compared are the same object, and **False** otherwise.
 - To check for **equality** or **inequality**, try:
tf.equal() and **tf.not_equal**, respectively.

TensorFlow graphs

- TensorFlow **automatically creates a Graph** when the library is loaded and assigns it to be the **default**.
- Thus, any Operations, tensors, etc. defined outside of a **Graph.as_default()** context manager will automatically be placed in the default graph

TensorFlow graphs

- If you'd like to get a handle to the default graph, use the `tf.get_default_graph()` function
- In most TensorFlow programs, you will only ever deal with the default graph
- When defining multiple graphs in one file, it's better to either not use the default graph or immediately assign a handle to it
- This ensures that nodes are added to each graph in a uniform manner

TensorFlow graphs

- Additionally, it is possible to load in previously defined models from other **TensorFlow** scripts and assign them to **Graph** objects
- This can be done by using a combination of the **graph.as_graph_def()** and **tf.import_graph_def** functions
- Thus, a user can compute and use the output of several separate models in the same Python file.

TensorFlow sessions

- **Sessions**, are responsible for **graph execution**
- The **constructor** takes in three optional parameters:
 - **target** specifies the execution engine to use:
 - For most applications, this will be left at its default empty string value. When using sessions in a distributed setting, this parameter is used to connect to **tf.train.Server** instances
 - **graph** specifies the **Graph** object that will be launched in the **Session**.
 - The default value is **None**, which indicates that the current default graph should be used. When using multiple graphs, it's best to explicitly pass in the **Graph** you'd like to run (instead of creating the **Session** inside of a **with** block).
 - **config** allows users to specify options to configure the session, such as limiting the number of CPUs or GPUs to use, setting optimization parameters for graphs, and logging options.

TensorFlow sessions

- Once a **Session** is opened, you can use its primary method, **run()**, to calculate the value of a desired Tensor output
- **Session.run()** takes in one required parameter, **fetches**,
(as well as **three optional** parameters: **feed_dict** , **options** , and **run_metadata**)

TensorFlow sessions

- We can also pass in a list of graph elements
- When **fetches** is a list, the output of **run()** will be a list with values corresponding to the output of the requested elements.
- In this example, we ask for the values of **a** and **b**, in that order
- Since both **a** and **b** are tensors, we receive their values as output

TensorFlow sessions

- In addition using `fetches` to get `Tensor` outputs, you'll also see examples where we give `fetches` a direct handle to an `Operation` which is a useful side-effect when run
- An example of this is `tf.global_variables_initializer()`, which prepares `all TensorFlow Variable objects` to be used
- We still pass the Op as the `fetches` parameter, but the result of `Session.run()` will be `None`

TensorFlow sessions

- The parameter **feed_dict** is used to override **Tensor** values in the graph, and it expects a Python dictionary object as input.
- The **keys** in the dictionary **are handles to Tensor objects** that should be overridden, while the **values** can be **numbers, strings, lists, or NumPy arrays** (as described previously)
- The **values must be of the same type** (or able to be converted to the same type) as the Tensor key.

TensorFlow placeholder

- Adding Inputs with **placeholder** nodes
- To take values from the client and plug them into our graph we use what is called a “**placeholder**”
- **Placeholders**, act as if they are Tensor objects, but they **do not have their values specified when created**
- Instead, **they hold the place for a Tensor** that will be fed at runtime, hence **become input nodes**

TensorFlow placeholder

- Adding Inputs with **placeholder** nodes
- **tf.placeholder** takes in a required parameter **dtype**, as well as the optional parameter **shape**:
 - **dtype** specifies the data type of values that will be passed into the placeholder. This is required, in order to ensure that there will be **no type mismatch errors**
 - **shape** specifies what shape the fed Tensor will be. The default value of **shape** is **None**, which means a Tensor of any shape will be accepted

TensorFlow variables

- **Tensor** and **Operation** objects are **immutable**, but machine learning tasks, by their nature, need a way to save changing values over time
- This is accomplished in TensorFlow with **Variable objects**, which contain **mutable tensor values** that persist across multiple calls to **Session.run()**.
- You can create a **Variable** by using its constructor, **tf.Variable()**

TensorFlow variables

- **Variables** can be used in TensorFlow functions/Operations anywhere you might use a **Tensor**
- Its present value will be passed on to the **Operation** using it
- The initial value of Variables will often be large tensors of zeros, ones, or random values
- TensorFlow has a number of helper Ops, such as:
`tf.zeros()`, `tf.ones()`, `tf.random_normal()`, and `tf.random_uniform()`

TensorFlow variables

- Instead of using `tf.random_normal()`, you'll often see use of `tf.truncated_normal()` instead, as it doesn't create any values more than two standard deviations away from its mean
- This prevents the possibility of having one or two numbers be significantly different than the other values in the tensor:
- You can pass in these Operations as the initial values of Variables as you would a handwritten Tensor

TensorFlow variables

- Variable objects **live in the Graph** like most other TensorFlow objects, but their state is actually **managed by a Session**.
- Because of this, **Variables** have an extra step involved in order to use them:
 - you **must initialize** the **Variable** within a **Session**

TensorFlow **changing** variables

- If you'd only like to initialize a subset of Variables defined in the graph, you can use `tf.initialize_variables()`
- This takes in a list of Variables to be initialized
- In order to change the value of the **Variable**, you can use the `Variable.assign()` method, which gives the **Variable** the new value to be

* Note that `Variable.assign()` is an **Operation**, and must be run in a **Session** to take effect

TensorFlow **changing** variables

- For simple incrementing and decrementing of Variables, TensorFlow includes the **Variable.assign_add()** **Variable.assign_sub()** methods
- Because **Sessions** maintain **Variable** values separately, each **Session** can have its own current value for a **Variable** defined in a graph

TensorFlow **changing** variables

- If you'd like to reset your Variables to their starting value, simply call `tf.global_variables_initializer()` again
- Or you can use `tf.variables_initializer()` if you only want to reset a subset of them)

TensorFlow **trainable** variables

- **Optimizer** classes automatically train **machine learning** models
- That means that it will change values of **Variable** objects without explicitly asking to do so
- If there are **Variables** in your graph that should only be changed manually and not with an **Optimizer**, you need to set their trainable parameter to False when creating them

TensorFlow **name scopes**

- So far, we've only worked with small graphs containing a few nodes and small tensors, but **real world models** can contain **dozens or hundreds of nodes**, as well as **millions of parameters**.
- In order to manage this level of complexity, TensorFlow currently offers a mechanism **to help organize your graphs** with:
name scopes
- They are incredibly simple to use and provide great value when visualizing your graph with TensorBoard.

TensorFlow **name scopes**

- **Name scopes** allow you to group Operations into larger, named blocks.
- When you launch your graph with TensorBoard, each name scope will encapsulate its own Ops, making the visualization much more digestible.
- For basic name scope usage, simply add your Operations in a with **tf.name_scope(<name>)** block (... see next slide)

TensorFlow **name scopes**

- To see the result of these name scopes in **TensorBoard**, let's open up a **FileWriter** and write this graph to disk

TensorFlow **name scopes**

- Because the **tf.summary.FileWriter** exports the graph immediately, we can simply start up **TensorBoard** after running the code on the previous slide.
- Navigate to where you ran the previous script and start up **TensorBoard**
- This will start a TensorBoard server on your local computer at port 6006.
- Open up a browser and enter **localhost:6006** into the URL bar

TensorFlow **name scopes**

- You'll notice that the **add** and **mul Operations** we added to the graph aren't immediately visible- instead, we see their enclosing name scopes.
- You can expand the name scope boxes by clicking on the plus **+** icon in their upper right corner.
- Inside of each scope, you'll see the individual **Operations** you've added to the graph

TensorFlow **name scopes**


- This model has:
 - two scalar placeholder nodes as input
 - a TensorFlow **constant**
 - a middle chunk called “**Transformation**”, and
 - a final **output node** that uses `tf.maximum()` as its Operation.
- We can see this high-level overview inside of TensorBoard

TensorFlow **name scopes**

- Inside of the **Transformation** name scope are four more name scopes arranged in two “**layers**”.
- The **first layer** is comprised of scopes “A” and “B”, which pass their output values into the next layer of “C” and “D”.
- The **final node** then uses the outputs from this last layer as its input.
- If you expand the Transformation name scope in **TensorBoard**, you’ll get a look the one shown on next slide

TensorFlow

- Below is how we can reconstruct our graph from Lecture 2:



```
basic_graph.py
1 # Import the tensorflow library, and reference it as 'tf'
2 import tensorflow as tf
3
4 # Build our graph nodes, starting from the inputs
5 a = tf.constant(5, name="input_a")
6 b = tf.constant(3, name="input_b")
7 c = tf.multiply(a,b, name="mul_c")
8 d = tf.add(a,b, name="add_d")
9 e = tf.add(c,d, name="add_e")
10
11 # Open up a TensorFlow Session
12 sess = tf.Session()
13
14 # Execute our output node, using our Session
15 sess.run(e)
16
17 # Open a TensorFlow FileWriter to write our graph to disk
18 writer = tf.summary.FileWriter('./my_graph', sess.graph)
19
20 # Close our FileWriter and Session objects
21 writer.close()
22 sess.close()
```

TensorFlow

- Here is a new implementation with a new graph:

```
3 # First we need to import TensorFlow:
4 import tensorflow as tf
5
6 # Defining our single input node:
7 a = tf.constant([5,3], name="input_a")
8
9 # Defining node 'b':
10 b = tf.reduce_prod(a, name="prod_b")
11
12 # Defining the next two nodes in our graph:
13 c = tf.reduce_sum(a, name="sum_c")
14
15 # This last line defines the final node in our graph:
16 d = tf.add(b,c, name="add_d")
17
18 # To run we have to add the two extra lines or run them in the shell:
19 sess = tf.Session()
20 sess.run(d)
```

TensorFlow


- We made few main changes here:
 1. We **replaced** the separate nodes **a** and **b** with a consolidated input node (**now just a**).
 2. We passed in a list of numbers, which **tf.constant** is able to **convert to a 1-D Tensor**
 3. Our **multiplication** and **addition** Operations, which used to take in scalar values, are now **tf.reduce_prod()** and **tf.reduce_sum()**
 4. These functions, when just given a Tensor as input, take all of its values and either multiply or sum them up, respectively

Data Science


Data Mining and Machine Learning concepts ...

Data Science


The Fields of Data Science



The Big Picture



The Big Picture



Data Mining

- What is *data mining*?
 - Data mining is defined as the process of discovering patterns in data
 - Data mining is about solving problems by analyzing data already present in datasets / databases
 - In data mining, the data is stored electronically and the search is automated
 - It has been estimated that the amount of data stored in the world's databases doubles every 20 months

Data Mining

- What is *data mining*?
 - Data mining is a topic that involves **learning in a practical**, nontheoretical sense
 - It is about finding and **describing previously unknown patterns** in data
 - The output may include a description of a structure that can be used to **classify unknown examples**
 - Finally it is about the **acquisition of knowledge** and the ability to use it

Data Mining

- Finding patterns in data that provide insight or enable fast and accurate decision making
- Strong, accurate patterns are needed to make decisions
 - Problem 1: most patterns are not interesting
 - Problem 2: patterns may be inexact (or spurious)
 - Problem 3: data may be garbled or missing
- Machine learning techniques identify patterns in data and provide many tools for data mining
- Of primary interest are machine learning techniques that provide structural descriptions